

# Week 07: Graph Algorithms 1

## Graph Algorithms

### Problems on Graphs

2/79

What kind of problems do we want to solve on/via graphs?

- is the graph fully-connected?
- can we remove an edge and keep it fully-connected?
- is one vertex reachable starting from some other vertex?
- what is the cheapest cost path from  $v$  to  $w$ ?
- which vertices are reachable from  $v$ ? (transitive closure)
- is there a cycle that passes through all vertices? (circuit)
- is there a tree that links all vertices? (spanning tree)
- what is the minimum spanning tree?
- ...
- can a graph be drawn in a plane with no crossing edges? (planar graphs)
- are two graphs "equivalent"? (isomorphism)

## Graph Algorithms

3/79

In this course we examine algorithms for

- connectivity (simple graphs)
- path finding (simple/directed graphs)
- minimum spanning trees (weighted graphs)
- shortest path (weighted graphs)

and look at generic methods for traversing graphs.

We begin with one of the simplest graph traversals ...

## Graph Traversal

### Finding a Path

5/79

Questions on paths:

- is there a path between two given vertices ( $src, dest$ )?
- what is the sequence of vertices from  $src$  to  $dest$ ?

Approach to solving problem:

- examine vertices adjacent to  $src$
- if any of them is  $dest$ , then done
- otherwise try vertices two edges from  $v$

- repeat looking further and further from  $v$


Two strategies for graph traversal/search: *depth-first*, *breadth-first*

- DFS follows one path to completion before considering others
- BFS "fans-out" from the starting vertex ("spreading" subgraph)

### ... Finding a Path

6/79

Comparison of BFS/DFS search for checking if there is a path from  $a$  to  $h$  ...



Both approaches ignore some edges by remembering previously visited vertices.

### ... Finding a Path

7/79

Depth-first ...

- favour following path rather than neighbours
- can be implemented recursively or iteratively (via stack)
- full traversal produces a *depth-first spanning tree*

Breadth-first ...

- favour neighbours rather than path following
- can be implemented iteratively (via queue)
- full traversal produces a *breadth-first spanning tree*


### Exercise #1: Traversal-induced Spanning Trees

8/79

A *spanning tree* of a graph


- includes all vertices
- uses a subset of edges, without cycles

Show the DFS and BFS spanning trees for the graph below, starting with 0:




Consider neighbours in ascending order

Answer:



DFS spanning tree



BFS spanning tree

## Depth-first Search

10/79

Depth-first search can be described recursively as

**depthFirst(G,v):**

1. mark  $v$  as visited
2. for each  $(v,w) \in \text{edges}(G)$  do
  - if  $w$  has not been visited then
  - depthFirst(w)**

The recursion induces *backtracking*

### ... Depth-first Search

11/79

Recursive DFS path checking

**hasPath(G,src,dest):**

**Input** graph  $G$ , vertices  $\text{src}, \text{dest}$   
**Output** true if there is a path from  $\text{src}$  to  $\text{dest}$  in  $G$ ,  
 false otherwise

**return** **dfsPathCheck(G,src,dest)**

**dfsPathCheck(G,v,dest):**

**mark**  $v$  **as visited**  
**for all**  $(v,w) \in \text{edges}(G)$  **do**  
 | **if**  $w = \text{dest}$  **then** // found edge to dest  
 | | **return** true  
 | **else if**  $w$  **has not been visited** **then**

```


    if dfsPathCheck(G,w,dest) then
        return true // found path via w to dest
    end if
  end if
end for
return false // no path from v to dest

```

### Exercise #2: Depth-first Traversal (i)

12/79

Trace the execution of **dfsPathCheck(G,0,5)** on:



Consider neighbours in ascending order

Answer:

0 - 1 - 2 - 3 - 4 - 5

### ... Depth-first Search

14/79

Cost analysis:

- each vertex visited at most once  $\Rightarrow$  cost =  $O(V)$
- visit all edges incident on visited vertices  $\Rightarrow$  cost =  $O(E)$ 
  - assuming an adjacency list representation

*Time complexity of DFS:*  $O(V+E)$  (adjacency list representation)

### ... Depth-first Search

15/79

Knowing whether a path exists can be useful

Knowing what the path is even more useful

$\Rightarrow$  record the previously visited node as we search through the graph (so that we can then trace path through graph)

Make use of global variable:

- **visited[ ]** ... array to store previously visited node, for each node being visited

visited[] // store previously visited node, for each vertex 0..nV-1


```
findPath(G,src,dest):
  Input graph G, vertices src,dest

  for all vertices v∈G do
    visited[v]=-1
  end for
  visited[src]=src                // starting node of the path
  if dfsPathCheck(G,src,dest) then // show path in dest..src order
    v=dest
    while v≠src do
      print v "-"
      v=visited[v]
    end while
    print src
  end if
```

```
dfsPathCheck(G,v,dest):
  for all (v,w)∈edges(G) do
    if visited[w]=-1 then
      visited[w]=v
      if w=dest then                // found edge from v to dest
        return true
      else if dfsPathCheck(G,w,dest) then
        return true                // found path via w to dest
      end if
    end if
  end for
  return false                      // no path from v to dest
```

Exercise #3: Depth-first Traversal (ii)

Show the DFS order in which we visit vertices in this graph when searching for a path from 0 to 6:



Consider neighbours in ascending order

0	0	3	5	3	1	5	4	7	8
[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]	[8]	[9]

Path: 6-5-1-0

... Depth-first Search

DFS can also be described non-recursively (via a *stack*):

visited[] // array of visiting orders, indexed by vertex 0..nV-1

```
findPathDFS(G,src,dest):
  Input graph G, vertices src,dest


  for all vertices v∈G do
    visited[v]=-1
  end for
  found=false
  visited[src]=src
  push src onto new stack s
  while ¬found ∧ s is not empty do
    pop v from s
    for each (v,w)∈edges(G) do
      if visited[w]=-1 then
        visited[w]=v
        if w=dest then
          found=true
        else
          push w onto s
        end if
      end if
    end for
  end while
  if found then
    display path in dest..src order
  end if
```

Uses standard stack operations (push, pop, check if empty)

Time complexity is the same:  $O(V+E)$  (each vertex added to stack once, each element in vertex's adjacency list visited once)

Exercise #4: Depth-first Traversal (iii)

Show how the stack evolves when executing findPathDFS(*g*,0,5) on:



Push neighbours in *descending* order ... so they get popped in ascending order

1        2        3

4        4        4        4

(empty) → 0 → 5 → 5 → 5 → 5 → 5 → (empty)

## Breadth-first Search

22/79

Basic approach to breadth-first search (BFS):

- visit and mark current vertex
- visit all neighbours of current vertex
- then consider neighbours of neighbours

Notes:

- tricky to describe recursively
- a minor variation on non-recursive DFS search works  
⇒ switch the *stack* for a *queue*

### ... Breadth-first Search

23/79

BFS algorithm (records visiting order):

visited[] // array of visiting orders, indexed by vertex 0..nV-1

```
findPathBFS(G,src,dest):
  Input  graph G, vertices src,dest

  for all vertices v∈G do
    visited[v]=-1
  end for
  found=false
  visited[src]=src
  enqueue src into new queue q
  while ¬found ∧ q is not empty do
    dequeue v from q
```


```
for each neighbour w of v do
  if visited[w]==-1 then
    visited[w]=v
    if w=dest then
      found=true
    else
      enqueue w into q
    end if
  end if
end for
end while
if found then
  display path in dest..src order
end if
```

Uses standard queue operations (enqueue, dequeue, check if empty)

### Exercise #5: Breadth-first Traversal

24/79

Show the BFS order in which we visit vertices in this graph when searching for a path from 0 to 6:



Consider neighbours in ascending order

0	0	0	2	5	0	5	5	-1	-1
[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]	[8]	[9]

Path: 6-5-0

### ... Breadth-first Search

26/79

Time complexity of BFS:  $O(V+E)$  (same as DFS)

BFS finds a "shortest" path

- based on minimum # edges between *src* and *dest*.
- stops with first-found path, if there are multiple ones

In many applications, edges are weighted and we want path

- based on minimum sum-of-weights along path  $src \dots dest$


We discuss weighted/directed graphs later.

## Other DFS Examples

27/79

Other problems to solve via DFS graph search

- checking for the existence of a cycle
- determining which connected component each vertex is in



## Exercise #6: Buggy Cycle Check

28/79

A graph has a *cycle* if

- it has a path of length  $> 1$
- with start vertex  $src$  = end vertex  $dest$
- and without using any edge more than once

We are not required to give the path, just indicate its presence.

The following DFS cycle check has two bugs. Find them.

```
hasCycle(G):
|   Input  graph G
|   Output true if G has a cycle, false otherwise
|
|   choose any vertex  $v \in G$ 
|   return dfsCycleCheck(G,v)
|
dfsCycleCheck(G,v):
|   mark v as visited
|   for each  $(v,w) \in \text{edges}(G)$  do
|   |   if w has been visited then    // found cycle
|   |   |   return true
|   |   else if dfsCycleCheck(G,w) then
|   |   |   return true
|   |   end if
|   end for
|   return false                // no cycle at v
```

1. Only one connected component is checked.
2. The loop

**for each**  $(v,w) \in \text{edges}(G)$  **do**

should exclude the neighbour of  $v$  from which you just came, so as to prevent a single edge  $w-v$  from being classified as a cycle.

## Computing Connected Components

30/79

Problems:

- how many connected subgraphs are there?
- are two vertices in the same connected subgraph?

Both of the above can be solved if we can

- build an array, one element for each vertex  $V$
- indicating which connected component  $V$  is in
- componentOf [ ] ... array  $[0..nV-1]$  of component IDs

## ... Computing Connected Components

31/79

Algorithm to assign vertices to connected components:


```
components(G):
|   Input  graph G
|
|   for all vertices  $v \in G$  do
|   |   componentOf[v] = -1
|   end for
|   compID = 0
|   for all vertices  $v \in G$  do
|   |   if componentOf[v] = -1 then
|   |   |   dfsComponents(G,v,compID)
|   |   |   compID = compID + 1
|   |   end if
|   end for
|
dfsComponents(G,v,id):
|   componentOf[v] = id
|   for all vertices w adjacent to v do
|   |   if componentOf[w] = -1 then
|   |   |   dfsComponents(G,w,id)
|   |   end if
|   end for
```

## Exercise #7: Connected components

32/79

Trace the execution of the algorithm

- 1. on the graph shown below
- 2. on the same graph but with the dotted edges added



Consider neighbours in ascending order

1.

[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]
-1	-1	-1	-1	-1	-1	-1	-1
0	-1	-1	-1	-1	-1	-1	-1
0	-1	0	-1	-1	-1	-1	-1
0	0	0	-1	-1	-1	-1	-1
0	0	0	1	-1	-1	-1	-1
...							
0	0	0	1	1	2	2	2

2.

[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]
-1	-1	-1	-1	-1	-1	-1	-1
0	-1	-1	-1	-1	-1	-1	-1
0	0	-1	-1	-1	-1	-1	-1
0	0	0	-1	-1	-1	-1	-1
0	0	0	0	-1	-1	-1	-1
...							
0	0	0	0	0	1	1	1

... Computing Connected Components

Consider an application where connectivity is critical

- we frequently ask questions of the kind above
- but we cannot afford to run components() each time

Add a new fields to the GraphRep structure:

```
typedef struct GraphRep *Graph;

struct GraphRep {
    ...
    int nC; // # connected components
    int *cc; // which component each vertex is contained in
    ... // i.e. array [0..nV-1] of 0..nC-1
}
```

... Computing Connected Components

With this structure, the above tasks become trivial:

```
// How many connected subgraphs are there?
int nConnected(Graph g) {
    return g->nC;
}

// Are two vertices in the same connected subgraph?
bool inSameComponent(Graph g, Vertex v, Vertex w) {
    return (g->cc[v] == g->cc[w]);
}
```

Consider maintenance of such a graph representation:

- initially, nC = nV (because no edges)
- adding an edge may reduce nC
- removing an edge may increase nC
- cc[] can simplify path checking (ensure v,w are in same component before starting search)

Additional maintenance cost amortised by reduced cost for nConnected() and inSameComponent()

Hamiltonian and Euler Paths

Hamiltonian Path and Circuit

Hamiltonian path problem:

- find a simple path connecting two vertices v,w in graph G
- such that the path includes each vertex exactly once

If v = w, then we have a Hamiltonian circuit

Simple to state, but difficult to solve (NP-complete)

Many real-world applications require you to visit all vertices of a graph:

- Travelling salesman


- Bus routes
- ...

Problem named after Irish mathematician, physicist and astronomer Sir William Rowan Hamilton (1805 - 1865)

### ... Hamiltonian Path and Circuit

38/79

Graph and two possible Hamiltonian paths:



### ... Hamiltonian Path and Circuit

39/79

Approach:

- generate all possible simple paths (using e.g. DFS)
- keep a counter of vertices visited in current path
- stop when find a path containing  $V$  vertices

Can be expressed via a recursive DFS algorithm

- similar to simple path finding approach, except
  - keeps track of path length; succeeds if length =  $v$
  - resets "visited" marker after unsuccessful path

### ... Hamiltonian Path and Circuit

40/79

Algorithm for finding Hamiltonian path:

visited[] // array [0.. $nV-1$ ] to keep track of visited vertices

```
hasHamiltonianPath(G,src,dest):
|   for all vertices v∈G do
|       visited[v]=false
|   end for
|   return hamiltonR(G,src,dest,#vertices(G)-1)
```


```
hamiltonR(G,v,dest,d):
|   Input G    graph
|           v    current vertex considered
|       dest    destination vertex
|           d     distance "remaining" until path found
|
|   if v=dest then
```

```
    if d=0 then return true else return false
else
|   visited[v]=true
|   for each (v,w)∈edges(G) ∧ ¬visited[w] do
|       if hamiltonR(G,w,dest,d-1) then
|           return true
|       end if
|   end for
end if
visited[v]=false           // reset visited mark
return false
```

### Exercise #8: Hamiltonian Path

41/79

Trace the execution of the algorithm when searching for a Hamiltonian path from 1 to 6:



Consider neighbours in ascending order

1-0-2-3-4-5-6	d≠0
1-0-2-3-4-5-7-8-9	no unvisited neighbour
1-0-2-3-4-5-7-9-8	no unvisited neighbour
1-0-2-3-4-7-5-6	d≠0
1-0-2-3-4-7-8-9	no unvisited neighbour
1-0-2-3-4-7-9-8	no unvisited neighbour
1-0-2-3-4-8-7-5-6	d≠0
1-0-2-3-4-8-7-9	no unvisited neighbour
1-0-2-3-4-8-9-7-5-6	✓


Repeat on your own with src=0 and dest=6

### ... Hamiltonian Path and Circuit

43/79

Analysis: worst case requires  $(V-1)!$  paths to be examined

Consider a graph with isolated vertex and the rest fully-connected



Checking `hasHamiltonianPath(g, 0, x)` for any  $x$

- requires us to consider every possible path
- e.g 1-2-3-4, 1-2-4-3, 1-3-2-4, 1-3-4-2, 1-4-2-3, ...
- starting from any  $x$ , there are  $3!$  paths  $\Rightarrow 4!$  total paths
- there is no path of length 5 in these  $(V-1)!$  possibilities

There is no known simpler algorithm for this task  $\Rightarrow$  *NP*-hard.

Note, however, that the above case could be solved in constant time if we had a fast check for 0 and  $x$  being in the same connected component


## Euler Path and Circuit

44/79


*Euler path* problem:

- find a path connecting two vertices  $v, w$  in graph  $G$
- such that the path includes each *edge* exactly once  
(note: the path does not have to be simple  $\Rightarrow$  can visit vertices more than once)

If  $v = w$ , then we have an *Euler circuit*



Euler Path: 4-3-1-5-2-1



Euler Circuit: 1-2-5-4-3-1

Many real-world applications require you to visit all edges of a graph:


- Postman
- Garbage pickup
- ...

## ... Euler Path and Circuit

45/79

Problem named after Swiss mathematician, physicist, astronomer, logician and engineer Leonhard Euler (1707 - 1783)

Based on a circuitous route via bridges in Königsberg





## Exercise #9: Euler Path

46/79

Is there a way to cross all the bridges of Königsberg exactly once on a walk through the town?

- treat land as nodes; bridges as edges



No

## ... Euler Path and Circuit

48/79

One possible "brute-force" approach:

- check for each path if it's an Euler path
- would result in factorial time performance

Can develop a better algorithm by exploiting:

*Theorem.* A graph has an Euler circuit if and only if it is connected and all vertices have even degree

*Theorem.* A graph has a non-circuitous Euler path if and only if it is connected and exactly two vertices have odd degree



## Exercise #10: Eulerian Graphs

49/79

Graphs with an Euler path are often called Eulerian Graphs

Which of these two graphs have an Euler path? an Euler circuit?





No Euler circuit

Only the second graph has an Euler path, e.g. 2-0-1-3-2-4-5

## ... Euler Path and Circuit

51/79

Assume the existence of  $\text{degree}(g, v)$  (degree of a vertex, cf. problem set week 6)

Algorithm to check whether a graph has an Euler path:

```
hasEulerPath(G,src,dest):
  Input  graph G, vertices src,dest
  Output true if G has Euler path from src to dest
         false otherwise

  if src≠dest then
    if degree(G,src) is even ∨ degree(G,dest) is even then
      return false
    end if
  else if degree(G,src) is odd then
    return false
  end if
  for all vertices v∈G do
    if v≠src ∧ v≠dest ∧ degree(G,v) is odd then
      return false
    end if
  end for
  return true
```

## ... Euler Path and Circuit

52/79

Analysis of hasEulerPath algorithm:

- assume that connectivity is already checked

- assume that degree is available via  $O(1)$  lookup
- single loop over all vertices  $\Rightarrow O(V)$

If degree requires iteration over vertices

- cost to compute degree of a single vertex is  $O(V)$
- overall cost is  $O(V^2)$

$\Rightarrow$  problem tractable, even for large graphs (unlike Hamiltonian path problem)

For the keen, a linear-time (in the number of edges,  $E$ ) algorithm to compute an Euler path is described in [Sedgewick] Ch.17.7.

## Directed Graphs

### Directed Graphs (Digraphs)

54/79

In our previous discussion of graphs:

- an edge indicates a relationship between two vertices
- an edge indicates nothing more than a relationship


In many real-world applications of graphs:

- edges are directional ( $v \rightarrow w \neq w \rightarrow v$ )
- edges have a *weight* (cost to go from  $v \rightarrow w$ )

### ... Directed Graphs (Digraphs)

55/79

Example digraph and adjacency matrix representation:



Undirectional  $\Rightarrow$  symmetric matrix

Directional  $\Rightarrow$  non-symmetric matrix

Maximum #edges in a digraph with  $V$  vertices:  $V^2$

### ... Directed Graphs (Digraphs)

56/79

Terminology for digraphs ...

*Directed path*: sequence of  $n \geq 2$  vertices  $v_1 \rightarrow v_2 \rightarrow \dots \rightarrow v_n$

- where  $(v_i, v_{i+1}) \in edges(G)$  for all  $v_i, v_{i+1}$  in sequence
- if  $v_1 = v_n$ , we have a *directed cycle*

*Degree of vertex:*  $deg(v)$  = number of edges of the form  $(v, \_ ) \in edges(G)$

- *Indegree* of vertex:  $deg^{-1}(v)$  = number of edges of the form  $(\_, v) \in edges(G)$

*Reachability:*  $w$  is reachable from  $v$  if  $\exists$  directed path  $v, \dots, w$

*Strong connectivity:* every vertex is reachable from every other vertex

*Directed acyclic graph* (DAG): graph containing no directed cycles

## Digraph Applications

57/79

Potential application areas:

Domain	Vertex	Edge
Web	web page	hyperlink
scheduling	task	precedence
chess	board position	legal move
science	journal article	citation
dynamic data	malloc'd object	pointer
programs	function	function call
make	file	dependency

### ... Digraph Applications

58/79

Problems to solve on digraphs:

- is there a directed path from  $s$  to  $t$ ? (transitive closure)
- what is the shortest path from  $s$  to  $t$ ? (shortest path)
- are all vertices mutually reachable? (strong connectivity)
- how to organise a set of tasks? (topological sort)
- which web pages are "important"? (PageRank)
- how to build a web crawler? (graph traversal)


## Digraph Representation

59/79

Similar set of choices as for undirectional graphs:

- array of edges (directed)
- vertex-indexed adjacency matrix (non-symmetric)
- vertex-indexed adjacency lists

$V$  vertices identified by  $0 \dots V-1$



### ... Digraph Representation

60/79

Costs of representations: (where degree  $deg(v)$  = #edges leaving  $v$ )

	array of edges	adjacency matrix	adjacency list
space usage	$E$	$V^2$	$V+E$
insert edge	$E$	$I$	$I$
exists edge $(v,w)$ ?	$E$	$I$	$deg(v)$
get edges leaving $v$	$E$	$V$	$deg(v)$

Overall, adjacency list representation is best

- real graphs tend to be sparse  
(large number of vertices, small average degree  $deg(v)$ )
- algorithms frequently iterate over edges from  $v$

## Reachability

62/79

## Transitive Closure

Given a digraph  $G$  it is potentially useful to know

- is vertex  $t$  reachable from vertex  $s$ ?

Example applications:

- can I complete a schedule from the current state?
- is a malloc'd object being referenced by any pointer?

How to compute transitive closure?

### ... Transitive Closure

63/79

One possibility:

- implement it via `hasPath(G,s,t)` (itself implemented by DFS or BFS algorithm)
- feasible if `reachable(G,s,t)` is infrequent operation

What if we have an algorithm that frequently needs to check reachability?

Would be very convenient/efficient to have:


```
reachable(G,s,t):
|   return G.tc[s][t]    // transitive closure matrix
```

Of course, if  $V$  is very large, then this is not feasible.


Exercise #11: Transitive Closure Matrix

64/79

Which reachable  $s \dots t$  exist in the following graph?



Transitive closure of example graph:



... Transitive Closure

66/79

**Goal:** produce a matrix of reachability values

- if  $tc[s][t]$  is 1, then  $t$  is reachable from  $s$
- if  $tc[s][t]$  is 0, then  $t$  is not reachable from  $s$

So, how to create this matrix?

Observation:

$\forall i,s,t \in \text{vertices}(G):$   
 $(s,i) \in \text{edges}(G) \wedge (i,t) \in \text{edges}(G) \Rightarrow tc[s][t] = 1$

$tc[s][t]=1$  if there is a path from  $s$  to  $t$  of length 2 ( $s \rightarrow i \rightarrow t$ )

67/79

... Transitive Closure

If we implement the above as:

```
make tc[][] a copy of edges[][]
for all ivertices(G) do
  for all svertices(G) do
    for all tvertices(G) do
      if tc[s][i]=1 ^ tc[i][t]=1 then
        tc[s][t]=1
      end if
    end for
  end for
end for
```

then we get an algorithm to convert edges into a  $tc$

This is known as *Warshall's algorithm*

68/79

... Transitive Closure

How it works ...

After iteration 1,  $tc[s][t]$  is 1 if

- either  $s \rightarrow t$  exists or  $s \rightarrow 0 \rightarrow t$  exists

After iteration 2,  $tc[s][t]$  is 1 if any of the following exist

- $s \rightarrow t$  or  $s \rightarrow 0 \rightarrow t$  or  $s \rightarrow 1 \rightarrow t$  or  $s \rightarrow 0 \rightarrow 1 \rightarrow t$  or  $s \rightarrow 1 \rightarrow 0 \rightarrow t$


Etc. ... so after the  $V^{th}$  iteration,  $tc[s][t]$  is 1 if

- there is any directed path in the graph from  $s$  to  $t$

Exercise #12: Transitive Closure

69/79

Trace Warshall's algorithm on the following graph:



1<sup>st</sup> iteration  $i=0$ :



tc	[0]	[1]	[2]	[3]
[0]	0	1	1	1
[1]	1	1	1	1
[2]	0	1	0	0
[3]	0	0	0	0

2<sup>nd</sup> iteration i=1:

tc	[0]	[1]	[2]	[3]
[0]	1	1	1	1
[1]	1	1	1	1
[2]	1	1	1	1
[3]	0	0	0	0

3<sup>rd</sup> iteration i=2: unchanged

4<sup>th</sup> iteration i=3: unchanged

## ... Transitive Closure

71/79

Cost analysis:

- storage: additional  $V^2$  items (each item may be 1 bit)
- computation of transitive closure:  $V^3$
- computation of `reachable()`:  $O(I)$  after having generated `tc[ ][ ]`

Amortization: would need many calls to `reachable()` to justify other costs

Alternative: use DFS in each call to `reachable()`

Cost analysis:

- storage: cost of queue and set during `reachable`
- computation of `reachable()`: cost of DFS =  $O(V^2)$  (for adjacency matrix)

## Digraph Traversal

72/79

Same algorithms as for undirected graphs:

**depthFirst(v):**

1. mark `v` as visited
2. for each  $(v, w) \in \text{edges}(G)$  do
  - if `w` has not been visited then
    - depthFirst(w)**

**breadth-first(v):**

1. enqueue `v`
2. while queue not empty do
  - dequeue `v`
  - if `v` not already visited then
    - mark `v` as visited
    - enqueue each vertex `w` adjacent to `v`

## Example: Web Crawling

73/79

**Goal:** visit every page on the web

**Solution:** breadth-first search with "implicit" graph

```
webCrawl(startingURL):
    mark startingURL as alreadySeen
    enqueue(Q, startingURL)
    while isEmpty(Q) do
        nextPage=dequeue(Q)
        visit nextPage
        for each hyperLink on nextPage do
            if hyperLink not alreadySeen then
                mark hyperLink as alreadySeen
                enqueue(Q, hyperLink)
            end if
        end for
    end while
```

`visit` scans page and collects e.g. keywords and links

## PageRank

74/79

**Goal:** determine which Web pages are "important"

**Approach:** ignore page contents; focus on hyperlinks

- treat Web as graph: page = vertex, hyperlink = di-edge
- pages with many incoming hyperlinks are important
- need to computing "incoming degree" for vertices

Problem: the Web is a *very* large graph

- approx.  $10^{14}$  pages,  $10^{15}$  hyperlinks

Assume for the moment that we could build a graph ...

Most frequent operation in algorithm "Does edge (v,w) exist?"

## ... PageRank

75/79

Simple PageRank algorithm:

```
PageRank(myPage):
|   rank=0
|   for each page in the Web do
|   |   if linkExists(page,myPage) then
|   |       rank=rank+1
|   |   end if
|   end for
|
Note: requires inbound link check (not outbound as assumed above for cost of representation)
```

... PageRank 76/79

$V$  = # pages in Web,  $E$  = # hyperlinks in Web

Costs for computing PageRank for each representation:

Representation	linkExists(v,w)	Cost
Adjacency matrix	edge[v][w]	$I$
Adjacency lists	inLL(list[v],w)	$\cong E/V$

Not feasible ...

- adjacency matrix ...  $V \cong 10^{14} \Rightarrow$  matrix has  $10^{28}$  cells
- adjacency list ...  $V$  lists, each with  $\cong 10$  hyperlinks  $\Rightarrow 10^{15}$  list nodes

So how to really do it?

... PageRank 77/79

Approach: the random web surfer

- if we randomly follow links in the web ...
- ... more likely to re-discover pages with many inbound links

```
curr=random page, prev=null
for a long time do
|   if curr not in array ranked[] then
|       rank[curr]=0
|   end if
|   rank[curr]=rank[curr]+1
|   if random(0,100)<85 then // with 85% chance ...
|       prev=curr
|       curr=choose hyperlink from curr // ... crawl on
|   else
|       curr=random page // avoid getting stuck
|       prev=null
|   end if
end for
```

Could be accomplished while we crawl web to build search index

Exercise #13: Implementing Facebook 78/79

Facebook could be considered as a giant "social graph"

- what are the vertices?
- what are the edges?
- are edges directional?

What kind of algorithm would ...

- help us find people that you might like to "befriend"?

Summary 79/79

- Graph traversal
  - depth-first search (DFS)
  - breadth-first search (BFS)
  - applications: path finding, connected components
- Hamiltonian paths/circuits, Euler paths/circuits
- Digraphs: representations, applications, reachability
  
- Suggested reading (Sedgewick):
  - Hamiltonian/Euler paths ... Ch.17.7
  - Graph search ... Ch.18.1-18.3,18.7
  - Digraphs ... Ch.19.1-19.3