

Weighted Graphs

Weighted Graphs

2/40

Graphs so far have considered

- edge = an association between two vertices/nodes
- may be a precedence in the association (directed)

Some applications require us to consider

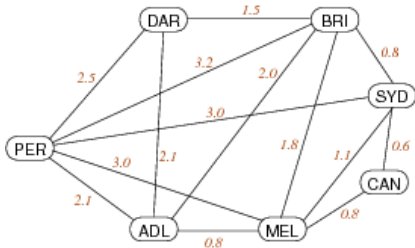
- a *cost* or *weight* of an association
- modelled by assigning values to edges (e.g. positive reals)

Weights can be used in both directed and undirected graphs.

... Weighted Graphs

3/40

Example: major airline flight routes in Australia



Representation: edge = direct flight; weight = approx flying time (hours)

... Weighted Graphs

4/40

Weights lead to minimisation-type questions, e.g.

1. Cheapest way to connect all vertices?
 - a.k.a. *minimum spanning tree* problem
 - assumes: edges are weighted and undirected
2. Cheapest way to get from *A* to *B*?
 - a.k.a. *shortest path* problem
 - assumes: edge weights positive, directed or undirected

Exercise #1: Implementing a Route Finder

5/40

If we represent a street map as a graph

- what are the vertices?
- what are the edges?
- are edges directional?
- what are the weights?
- are the weights fixed?

What kind of algorithm would ...

- help us find the "quickest" way to get from A to B?

Weighted Graph Representation

6/40

Weights can easily be added to:

- adjacency matrix representation (0/1 \rightarrow int or float)
- adjacency lists representation (add int/float to list node)

An alternative representation useful in this context:

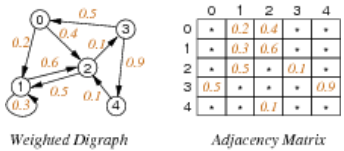
- edge list representation (list of (s,t,w) triples)

All representations work whether edges are directed or not.

... Weighted Graph Representation

7/40

Adjacency matrix representation with weights:

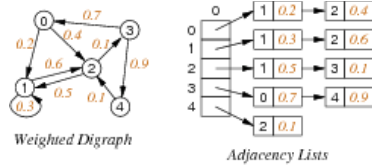


Note: need distinguished value to indicate "no edge".

... Weighted Graph Representation

8/40

Adjacency lists representation with weights:

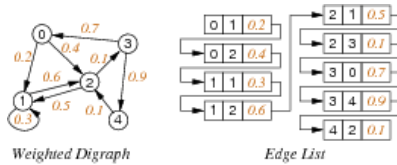


Note: if undirected, each edge appears twice with same weight

... Weighted Graph Representation

9/40

Edge array / edge list representation with weights:



Note: not very efficient for use in processing algorithms, but does give a possible representation for min spanning trees or shortest paths

... Weighted Graph Representation

10/40

Sample adjacency matrix implementation in C requires minimal changes to previous Graph ADT:

WGraph.h

```
// edges are pairs of vertices (end-points) plus positive weight
typedef struct Edge {
    Vertex v;
    Vertex w;
    int weight;
} Edge;
```

```
// returns weight, or 0 if vertices not adjacent
int adjacent(Graph, Vertex, Vertex);
```

... Weighted Graph Representation

11/40

WGraph.c

```
typedef struct GraphRep {
    int **edges; // adjacency matrix storing positive weights
                // 0 if nodes not adjacent
    int nV;      // #vertices
    int nE;      // #edges
} GraphRep;
```

```
void insertEdge(Graph g, Edge e) {
    assert(g != NULL && validV(g,e.v) && validV(g,e.w));
    if (g->edges[e.v][e.w] == 0) { // edge e not in graph
        g->edges[e.v][e.w] = e.weight;
        g->edges[e.w][e.v] = e.weight;
        g->nE++;
    }
}
```

```
int adjacent(Graph g, Vertex v, Vertex w) {
    assert(g != NULL && validV(g,v) && validV(g,w));
    return g->edges[v][w];
}
```

Minimum Spanning Trees

Minimum Spanning Trees

13/40

Reminder: *Spanning tree* ST of graph $G(V,E)$

- *spanning* = all vertices, *tree* = no cycles
- ST is a subgraph of G ($G'=(V,E')$ where $E' \subseteq E$)
- ST is *connected* and *acyclic*

Minimum spanning tree MST of graph G

- MST is a spanning tree of G
- sum of edge weights is no larger than any other ST

Applications: Computer networks, Electrical grids, Transportation networks ...

Problem: how to (efficiently) find MST for graph G ?

NB: MST may not be unique (e.g. all edges have same weight \Rightarrow every ST is MST)

... Minimum Spanning Trees

14/40

Brute force solution:

```
findMST(G):
|   Input graph G
|   Output a minimum spanning tree of G
|
|   bestCost=∞
|   for all spanning trees t of G do
|       if cost(t)<bestCost then
|           bestTree=t
|           bestCost=cost(t)
|   end if
```

```

|   end for
|   return bestTree

```

Example of *generate-and-test* algorithm.

Not useful because *#spanning trees* is potentially large (e.g. n^{n-2} for a complete graph with n vertices)

... Minimum Spanning Trees

Simplifying assumption:

- edges in *G* are not directed (MST for digraphs is harder)

Kruskal's Algorithm

One approach to computing MST for graph *G* with *V* nodes:

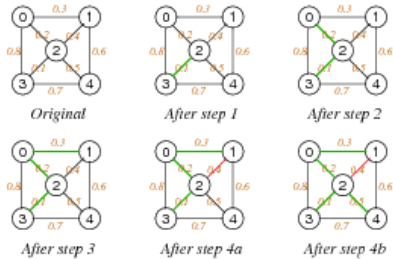
- start with empty MST
- consider edges in increasing weight order
 - add edge if it does not form a cycle in MST
- repeat until *V-1* edges are added

Critical operations:

- iterating over edges in weight order
- checking for cycles in a graph

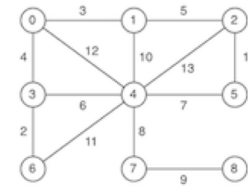
... Kruskal's Algorithm

Execution trace of Kruskal's algorithm:

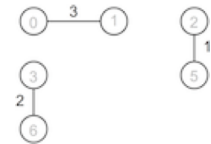


Exercise #2: Kruskal's Algorithm

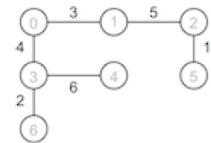
Show how Kruskal's algorithm produces an MST on:



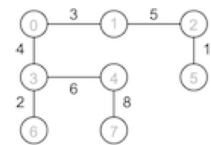
After 3rd iteration:



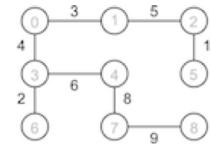
After 6th iteration:



After 7th iteration:



After 8th iteration (*V*-1=8 edges added):



... Kruskal's Algorithm

Pseudocode:

```

KruskalMST(G) :
|   Input  graph G with n nodes
|   Output a minimum spanning tree of G
|
|   MST=empty graph

```

```

sort edges(G) by weight
for each e in sortedEdgeList do
    MST = MST ∪ {e}
    if MST has a cycle then
        MST = MST \ {e}
    end if
    if MST has n-1 edges then
        return MST
    end if
end for

```

... Kruskal's Algorithm

21/40

Rough time complexity analysis ...

- sorting edge list is $O(E \cdot \log E)$
- at least V iterations over sorted edges
- on each iteration ...
 - getting next lowest cost edge is $O(1)$
 - checking whether adding it forms a cycle: cost = ??

Possibilities for cycle checking:

- use DFS ... too expensive?
- could use *Union-Find data structure* (see Sedgewick Ch.1)

Prim's Algorithm

22/40

Another approach to computing MST for graph $G=(V,E)$:

1. start from any vertex s and empty MST
2. choose edge not already in MST to add to MST
 - must be incident on a vertex already connected to s in MST
 - must have minimal weight of all such edges
3. repeat until MST covers all vertices

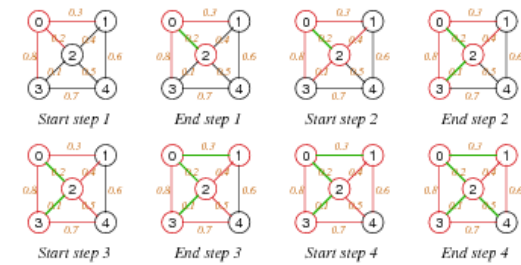
Critical operations:

- checking for vertex being connected in a graph
- finding min weight edge in a set of edges

... Prim's Algorithm

23/40

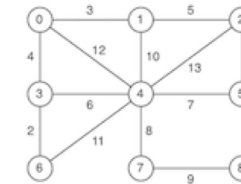
Execution trace of Prim's algorithm (starting at $s=0$):



Exercise #3: Prim's Algorithm

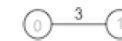
24/40

Show how Prim's algorithm produces an MST on:

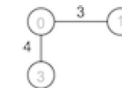


Start from vertex 0

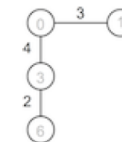
After 1st iteration:



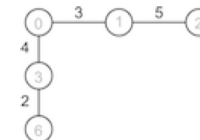
After 2nd iteration:



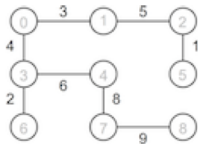
After 3rd iteration:



After 4th iteration:



After 8th iteration (all vertices covered):



... Prim's Algorithm

26/40

Pseudocode:

```
PrimMST(G):
  Input  graph G with n nodes
  Output a minimum spanning tree of G

  MST=empty graph
  usedV={0}
  unusedE=edges(g)
  while |usedV|<n do
    find e=(s,t,w)∈unusedE such that {
      s∈usedV ∧ t∉usedV ∧ w is min weight of all such edges
    }
    MST = MST ∪ {e}
    usedV = usedV ∪ {t}
    unusedE = unusedE \ {e}
  end while
  return MST
```

Critical operation: finding best edge

... Prim's Algorithm

27/40

Rough time complexity analysis ...

- V iterations of outer loop
- in each iteration ...
 - find min edge with set of edges is $O(E) \Rightarrow O(V \cdot E)$ overall
 - find min edge with *priority queue* is $O(\log E) \Rightarrow O(V \cdot \log E)$ overall

Note:

- Using a *priority queue* gives a variation of DFS (stack) and BFS (queue) graph traversal

Sidetrack: Priority Queues

28/40

Some applications of queues require

- items processed in order of "key"

- rather than in order of entry (FIFO — first in, first out)

Priority Queues (PQueues) provide this via:

- join: insert item into PQueue (replacing enqueue)
- leave: remove item with largest key (replacing dequeue)

... Sidetrack: Priority Queues

29/40

Comparison of different Priority Queue representations:

	sorted array	unsorted array	sorted list	unsorted list
space usage	$MaxN$	$MaxN$	$O(N)$	$O(N)$
join	$O(N)$	$O(1)$	$O(N)$	$O(1)$
leave	$O(N)$	$O(N)$	$O(1)$	$O(N)$
is empty?	$O(1)$	$O(1)$	$O(1)$	$O(1)$

for a PQueue containing N items

Other MST Algorithms

30/40

Boruvka's algorithm ... complexity $O(E \cdot \log V)$

- the oldest MST algorithm
- start with V separate components
- join components using min cost links
- continue until only a single component

Karger, Klein, and Tarjan ... complexity $O(E)$

- based on Boruvka, but non-deterministic
- randomly selects subset of edges to consider
- for the keen, here's [the paper](#) describing the algorithm

Shortest Path

32/40

Shortest Path

Path = sequence of edges in graph $G \quad p = (v_0, v_1), (v_1, v_2), \dots, (v_{m-1}, v_m)$

cost(path) = sum of edge weights along path

Shortest path between vertices s and t

- a simple path $p(s, t)$ where $s = first(p)$, $t = last(p)$

- no other simple path $q(s,t)$ has $cost(q) < cost(p)$

Assumptions: weighted digraph, no negative weights.

Finding shortest path between two given nodes known as *source-target* SPP problem

Variations: *single-source*, *all-pairs*

Applications: robot navigation, routing in data networks, ...

Single-source Shortest Path (SSSP)

33/40

Given: weighted digraph G , source vertex s

Result: shortest paths from s to all other vertices

- $dist[]$ V -indexed array of cost of shortest path from s
- $pred[]$ V -indexed array of predecessor in shortest path from s

Example:



Edge Relaxation

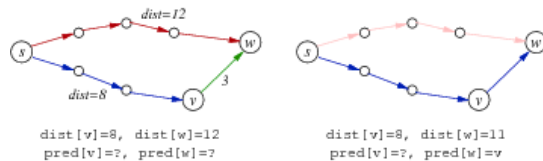
34/40

Assume: $dist[]$ and $pred[]$ as above (but containing data for shortest paths *discovered so far*)

$dist[v]$ is length of shortest known path from s to v

$dist[w]$ is length of shortest known path from s to w

Relaxation updates data for w if we find a shorter path from s to w :



Relaxation along edge $e=(v,w,weight)$:

- if $dist[v]+weight < dist[w]$ then
update $dist[w]:=dist[v]+weight$ and $pred[w]:=v$

35/40

Dijkstra's Algorithm

One approach to solving single-source shortest path ...

Data: $G, s, dist[], pred[]$ and

- $vSet$: set of vertices whose shortest path from s is known

Algorithm:

$dist[]$ // array of cost of shortest path from s
 $pred[]$ // array of predecessor in shortest path from s

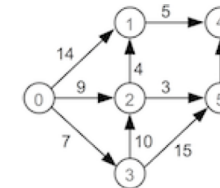
```
dijkstraSSSP(G,source):
    Input graph G, source node

    initialise dist[] to all ∞, except dist[source]=0
    initialise pred[] to all -1
    vSet=all vertices of G
    while vSet≠∅ do
        find s∈vSet with minimum dist[s]
        for each (s,t,w)∈edges(G) do
            relax along (s,t,w)
        end for
        vSet=vSet\{s}
    end while
```

Exercise #4: Dijkstra's Algorithm

36/40

Show how Dijkstra's algorithm runs on (source node = 0):



	[0]	[1]	[2]	[3]	[4]	[5]
dist	0	∞	∞	∞	∞	∞
pred	-	-	-	-	-	-

dist	0	14	9	7	∞	∞
pred	-	0	0	0	-	-

--	--	--	--	--	--	--

dist	0	14	9	7	∞	22
pred	–	0	0	0	–	3

dist	0	13	9	7	∞	12
pred	–	2	0	0	–	2

dist	0	13	9	7	20	12
pred	–	2	0	0	5	2

dist	0	13	9	7	18	12
pred	–	2	0	0	1	2

... Dijkstra's Algorithm

38/40

Why Dijkstra's algorithm is correct:

Hypothesis.

(a) For visited s ... $dist[s]$ is shortest distance from source

(b) For unvisited t ... $dist[t]$ is shortest distance from source *via visited nodes*

Proof.

Base case: no visited nodes, $dist[source]=0$, $dist[s]=\infty$ for all other nodes

Induction step:

1. If s is unvisited node with minimum $dist[s]$, then $dist[s]$ is shortest distance from source to s :
 - if \exists shorter path via only visited nodes, then $dist[s]$ would have been updated when processing the predecessor of s on this path
 - if \exists shorter path via an unvisited node u , then $dist[u] < dist[s]$, which is impossible if s has min distance of all unvisited nodes
2. This implies that (a) holds for s after processing s
3. (b) still holds for all unvisited nodes t after processing s :
 - if \exists shorter path via s we would have just updated $dist[t]$
 - if \exists shorter path without s we would have found it previously

... Dijkstra's Algorithm

39/40

Time complexity analysis ...

Each edge needs to be considered once $\Rightarrow O(E)$.

Outer loop has $O(V)$ iterations.

Implementing "**find** $s \in VSet$ **with** minimum $dist[s]$ "

1. try all $s \in VSet \Rightarrow cost = O(V) \Rightarrow overall\ cost = O(E + V^2) = O(V^2)$
2. using a PQueue to implement extracting minimum
 - can improve overall cost to $O(E + V \cdot \log V)$ (for best-known implementation)

Summary

- Weighted graph representations
- Minimum Spanning Tree (MST)
 - Kruskal, Prim
- Single source shortest path problem
 - Dijkstra
- Suggested reading:
 - Sedgewick, Ch.19.3,20-20.4,21-21.3

Produced: 11 Sep 2017