

Week 06

Things to Note ...

- Congratulations on finishing your first assignment!

In This Lecture ...

- Graph data structures ([Slides](#), [\[S\]](#) 17.1-17.5)

Coming Up ...

- Graph algorithms ([Slides](#), [\[S\]](#) Ch.18)

Nerds You Should Know

The next in a series on famous computer scientists ...



What he invented affects your life every single day ...

Nerds You Should Know (cont)

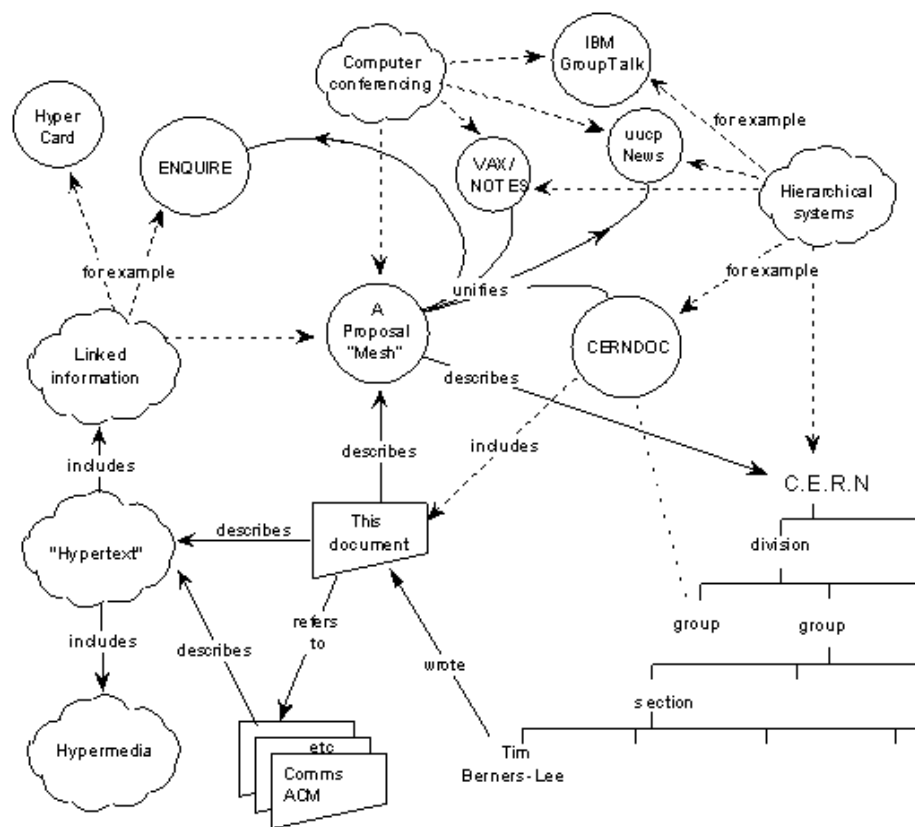
Sir Tim Berners-Lee



- Oxford CS Graduate (1976)
- Software engineer at CERN
- Founder/Director of W3C at MIT (1994)
- Inventing the Web ...
 - distributed hypertext
 - linking heterogeneous documents
 - universal naming scheme (URL)
 - transfer protocol (http)
 - later apologised for initial pair of slashes ('//') in a web address
 - also thinks he should have defined web addresses the other way round (`au.edu.unsw.cse`)
- Winner of the Turing Award in 2016

Nerds You Should Know (cont)

Tim Berners-Lee's original diagram of the "Web"



(from his proposal document, 1989)

Graph Definitions

Graphs

Many applications require

- a collection of **items** (i.e. a set)
- **relationships**/connections between items

Examples:

- maps: items are cities, connections are roads
- web: items are pages, connections are hyperlinks

Collection types you're familiar with

- lists ... linear sequence of items (week 3, COMP9021)
- trees ... branched hierarchy of items (COMP9021)

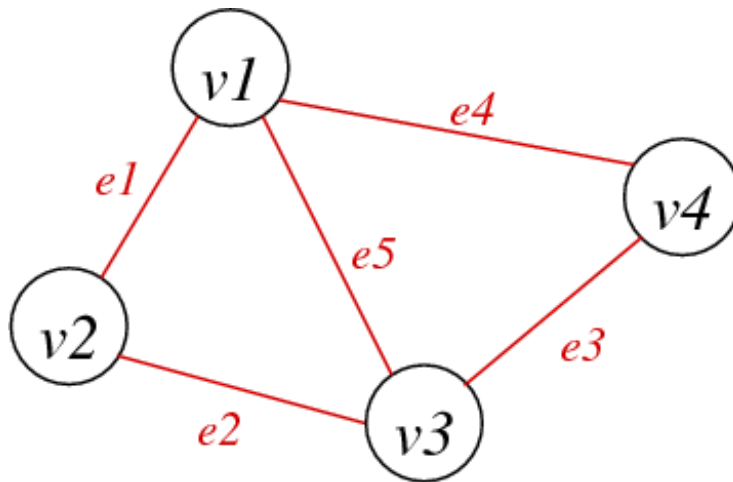
Graphs are more general ... allow arbitrary connections

Graphs (cont)

A graph $G = (V, E)$

- V is a set of vertices
- E is a set of edges (subset of $V \times V$)

Example:



$$V = \{v1, v2, v3, v4\}$$

$$E = \{e1, e2, e3, e4, e5\}$$

Graphs (cont)

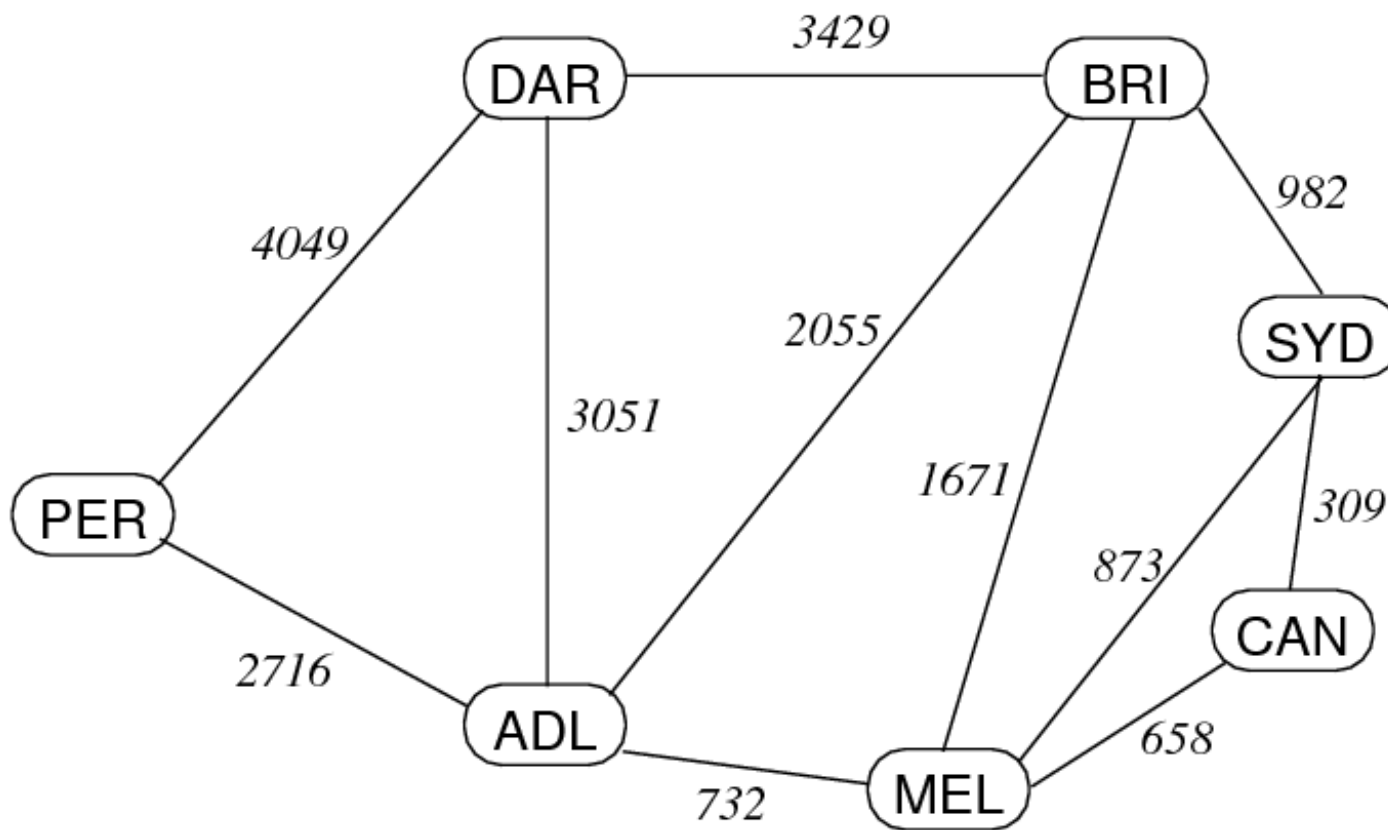
A real example: Australian road distances

Distance	Adelaide	Brisbane	Canberra	Darwin	Melbourne	Perth	Sydney
Adelaide	-	2055	1390	3051	732	2716	1605
Brisbane	2055	-	1291	3429	1671	4771	982
Canberra	1390	1291	-	4441	658	4106	309
Darwin	3051	3429	4441	-	3783	4049	4411
Melbourne	732	1671	658	3783	-	3448	873
Perth	2716	4771	4106	4049	3448	-	3972
Sydney	1605	982	309	4411	873	3972	-

Notes: vertices are cities, edges are distance between cities, symmetric

Graphs (cont)

Alternative representation of above:



Graphs (cont)

Questions we might ask about a graph:

- is there a way to get from item A to item B?
- what is the best way to get from A to B?
- which items are connected?

Graph algorithms are generally more complex than tree/list ones:

- no implicit order of items
- graphs may contain cycles
- concrete representation is less obvious
- algorithm complexity depends on connection complexity

Properties of Graphs

Terminology: $|V|$ and $|E|$ (cardinality) normally written just as V and E .

A graph with V vertices has at most $V(V-1)/2$ edges.

The ratio $E:V$ can vary considerably.

- if E is closer to V^2 , the graph is **dense**
- if E is closer to V , the graph is **sparse**
 - Example: web pages and hyperlinks

Knowing whether a graph is sparse or dense is important

- may affect choice of data structures to represent graph
- may affect choice of algorithms to process graph

Exercise #1: Number of Edges

The edges in a graph represent pairs of connected vertices. A graph with V has V^2 such pairs.

Consider $V = \{1, 2, 3, 4, 5\}$ with all possible pairs:

$$E = \{ (1, 1), (1, 2), (1, 3), (1, 4), (1, 5), (2, 1), (2, 2), \dots, (4, 5), (5, 5) \}$$

Why do we say that the maximum #edges is $V(V-1)/2$?

... because

- (v,w) and (w,v) denote the same edge (in an undirected graph)
- we do not consider loops (v,v)

Graph Terminology

For an edge e that connects vertices v and w

- v and w are **adjacent** (neighbours)
- e is **incident** on both v and w

Degree of a vertex v

- number of edges incident on v

Synonyms:

- vertex = node, edge = arc = link (Note: some people use arc for *directed* edges)

Graph Terminology (cont)

Path: a sequence of vertices where


- each vertex has an edge to its predecessor

Cycle: a path where


- last vertex in path is same as first vertex in path

Length of path or cycle:

- #edges



Path: 1-2, 2-3, 3-4



Cycle: 1-2, 2-3, 3-4, 4-1

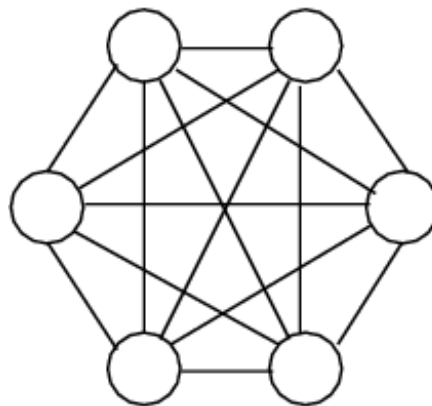
Graph Terminology (cont)

Connected graph

- there is a *path* from each vertex to every other vertex
- if a graph is not connected, it has ≥ 2 connected components

Complete graph K_V

- there is an *edge* from each vertex to every other vertex
- in a complete graph, $E = V(V-1)/2$



*Complete
Graph*

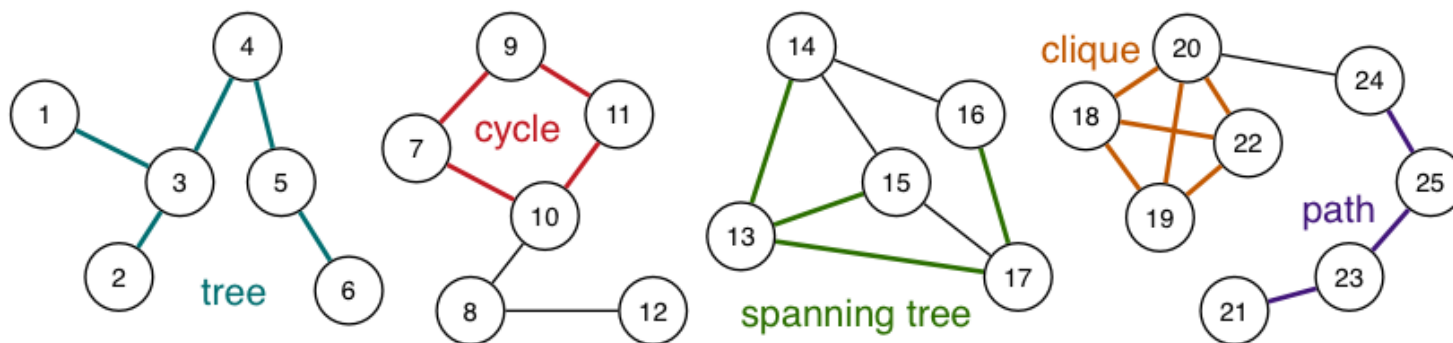
Graph Terminology (cont)

Tree: connected (sub)graph with no cycles

Spanning tree: tree containing all vertices

Clique: complete subgraph

Consider the following single graph:



This graph has 25 vertices, 32 edges, and 4 connected components

Note: The entire graph has no spanning tree; what is shown in green is a spanning tree of the third connected component

Graph Terminology (cont)


A **spanning tree** of connected graph $G = (V, E)$

- is a subgraph of G containing all of V
- and is a single tree (connected, no cycles)

A **spanning forest** of non-connected graph $G = (V, E)$

- is a subgraph of G containing all of V
- and is a set of trees (not connected, no cycles),
 - with one tree for each *connected component*

Exercise #2: Graph Terminology



$$V = \{v1, v2, v3, v4\}$$

$$E = \{e1, e2, e3, e4, e5\}$$

1. How many edges to remove to obtain a spanning tree?
2. How many different spanning trees?

1. 2

$$\frac{5 \cdot 4}{2}$$

2. $\frac{5 \cdot 4}{2} - 2 = 8$ spanning trees (no spanning tree if we remove $\{e1, e2\}$ or $\{e3, e4\}$)

Graph Terminology (cont)

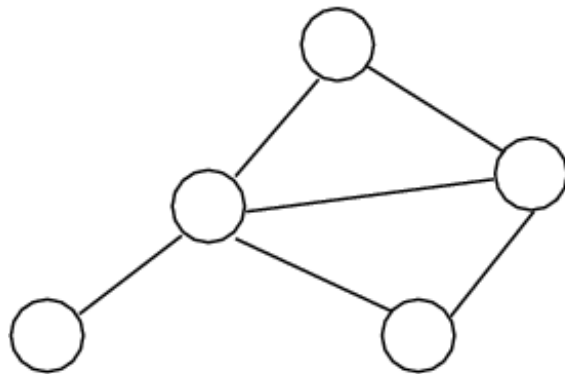
Undirected graph

- $edge(u,v) = edge(v,u)$, no self-loops (i.e. no $edge(v,v)$)

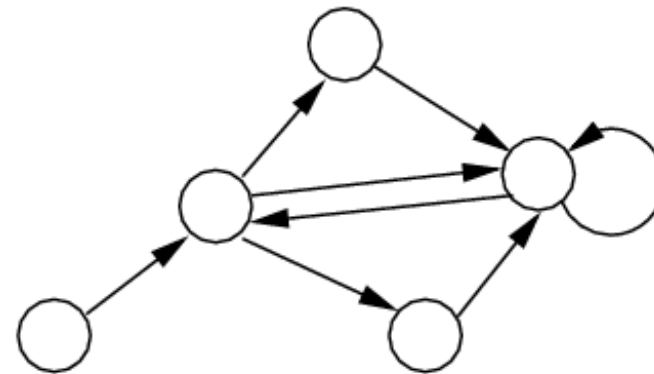
Directed graph

- $edge(u,v) \neq edge(v,u)$, can have self-loops (i.e. $edge(v,v)$)

Examples:



Undirected graph



Directed graph

Graph Terminology (cont)

Other types of graphs ...

Weighted graph

- each edge has an associated value (weight)
- e.g. road map (weights on edges are distances between cities)

Multi-graph

- allow multiple edges between two vertices
- e.g. function call graph (**f**() calls **g**() in several places)


Graph Data Structures

Graph Representations


Defining graphs:

- need some way of identifying vertices
- could give diagram showing edges and vertices
- could give a list of edges

E.g. four representations of the same graph:



(a)



(b)

1-2 1-3 1-4
2-4
3-4
4-5

(c)

1-3
2-1 2-4
4-1 4-3
5-4

(d)

Graph Representations (cont)

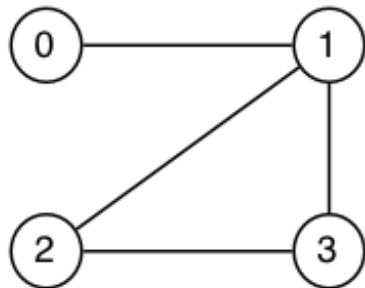
We will discuss three different graph data structures:

1. Array of edges
2. Adjacency matrix
3. Adjacency list

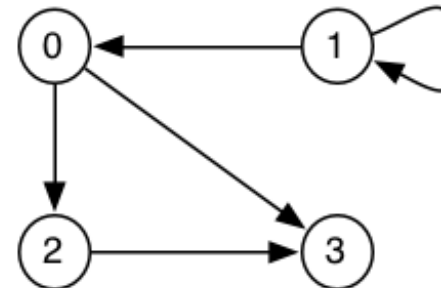
Array-of-edges Representation

Edges are represented as an array of **Edge** values (= pairs of vertices)

- space efficient representation
- adding and deleting edges is slightly complex
- undirected: order of vertices in an **Edge** doesn't matter
- directed: order of vertices in an **Edge** encodes direction



[(0,1), (1,2), (1,3), (2,3)]



[(1,0), (1,1), (0,2), (0,3), (2,3)]

For simplicity, we always assume vertices to be numbered $0 \dots v-1$

Array-of-edges Representation (cont)

Graph initialisation

```
newGraph(V):
    Input    number of nodes V
    Output  new empty graph

    g.nV = V    // #vertices (numbered 0..V-1)
    g.nE = 0    // #edges
    allocate enough memory for g.edges[]
    return g
```

Array-of-edges Representation (cont)

Edge insertion

```
insertEdge(g, (v,w)):  
    Input graph g, edge (v,w)  
  
    i=0  
    while i<g.nE  $\wedge$  (v,w) $\neq$ g.edges[i] do  
        i=i+1  
    end while  
    if i=g.nE then                                // (v,w) not found  
        g.edges[i]=(v,w)  
        g.nE=g.nE+1  
    end if
```

Array-of-edges Representation (cont)

Edge removal

```
removeEdge(g, (v,w)) :
    Input graph g, edge (v,w)

    i=0
    while i<g.nE  $\wedge$  (v,w) $\neq$ g.edges[i] do
        i=i+1
    end while
    if i<g.nE then                                // (v,w) found
        g.edges[i]=g.edges[g.nE-1] // replace by last edge in array
        g.nE=g.nE-1
    end if
```

Cost Analysis

Storage cost: $O(E)$

Cost of operations:

- initialisation: $O(1)$
- insert edge: $O(E)$ (assuming edge array has space)
- delete edge: $O(E)$ (need to find edge in edge array)

If array is full on insert

- allocate space for a bigger array, copy edges across \Rightarrow still $O(E)$

If we maintain edges in order

- use binary search to find edge $\Rightarrow O(\log E)$

Exercise #3: Array-of-edges Representation


Assuming an array-of-edges representation ...

Write an algorithm to output all edges of the graph

```
show(g):  
  Input graph g  
  
  for all i=0 to g.nE-1 do  
    print g.edges[i]  
  end for
```



Adjacency Matrix Representation

Edges represented by a $V \times V$ matrix



Undirected graph

A	0	1	2	3
0	0	1	0	1
1	1	0	0	1
2	0	0	0	1
3	1	1	1	0



Directed graph

A	0	1	2	3
0	0	0	0	1
1	1	0	0	1
2	0	0	0	0
3	0	0	1	0

Adjacency Matrix Representation (cont)

Advantages

- easily implemented as 2-dimensional array
- can represent graphs, digraphs and weighted graphs
 - graphs: symmetric boolean matrix
 - digraphs: non-symmetric boolean matrix
 - weighted: non-symmetric matrix of weight values

Disadvantages:

- if few edges (sparse) \Rightarrow memory-inefficient

Adjacency Matrix Representation (cont)

Graph initialisation

```
newGraph(V):  
  Input   number of nodes V  
  Output new empty graph  
  
  g.nV = V    // #vertices (numbered 0..V-1)  
  g.nE = 0    // #edges  
  allocate memory for g.edges[][]  
  for all i,j=0..V-1 do  
    g.edges[i][j]=0    // false  
  end for  
  return g
```

Adjacency Matrix Representation (cont)

Edge insertion

```
insertEdge(g, (v,w)) :  
    Input graph g, edge (v,w)  
  
    if g.edges[v][w]=0 then    // (v,w) not in graph  
        g.edges[v][w]=1        // set to true  
        g.edges[w][v]=1  
        g.nE=g.nE+1  
    end if
```

Adjacency Matrix Representation (cont)

Edge removal

```
removeEdge(g, (v,w)):  
    Input graph g, edge (v,w)  
  
    if g.edges[v][w]≠0 then    // (v,w) in graph  
        g.edges[v][w]=0        // set to false  
        g.edges[w][v]=0  
        g.nE=g.nE-1  
    end if
```

Exercise #4: Show Graph

Assuming an adjacency matrix representation ...

Write an algorithm to output all edges of the graph (no duplicates!)

Adjacency Matrix Representation (cont)

```
show(g):  
  Input graph g  
  
  for all i=0 to g.nV-1 do  
    for all j=i+1 to g.nV-1 do  
      if g.edges[i][j]≠0 then  
        print i-"-"j  
      end if  
    end for  
  end for
```

Exercise #5:

Analyse storage cost and time complexity of adjacency matrix representation

Storage cost: $O(V^2)$


If the graph is sparse, most storage is wasted.

Cost of operations:


- initialisation: $O(V^2)$ (initialise $V \times V$ matrix)
- insert edge: $O(1)$ (set two cells in matrix)
- delete edge: $O(1)$ (unset two cells in matrix)

Adjacency Matrix Representation (cont)

A storage optimisation: store only top-right part of matrix.



Undirected graph




New storage cost: $V-1$ int ptrs + $V(V+1)/2$ ints (but still $O(V^2)$)

Requires us to always use edges (v,w) such that $v < w$.

Adjacency List Representation

For each vertex, store linked list of adjacent vertices:




Undirected graph

$$A[0] = \langle 1, 3 \rangle$$

$$A[1] = \langle 0, 3 \rangle$$

$$A[2] = \langle 3 \rangle$$

$$A[3] = \langle 0, 1, 2 \rangle$$



Directed graph

$$A[0] = \langle 3 \rangle$$

$$A[1] = \langle 0, 3 \rangle$$

$$A[2] = \langle \rangle$$

$$A[3] = \langle 2 \rangle$$

Adjacency List Representation (cont)

Advantages

- relatively easy to implement in languages like C
- can represent graphs and digraphs
- memory efficient if $E:V$ relatively small

Disadvantages:

- one graph has many possible representations
(unless lists are ordered by same criterion e.g. ascending)

Adjacency List Representation (cont)

Graph initialisation

```
newGraph(V):  
  Input   number of nodes V  
  Output new empty graph  
  
  g.nV = V    // #vertices (numbered 0..V-1)  
  g.nE = 0    // #edges  
  allocate memory for g.edges[]  
  for all i=0..V-1 do  
    g.edges[i]=NULL    // empty list  
  end for  
  return g
```

Adjacency List Representation (cont)

Edge insertion:

```
insertEdge(g, (v,w)) :  
  Input graph g, edge (v,w)  
  
  if  $\neg$ inLL(g.edges[v],w) then // (v,w) not in graph  
    insertLL(g.edges[v],w)  
    insertLL(g.edges[w],v)  
    g.nE=g.nE+1  
  end if
```

Adjacency List Representation (cont)

Edge removal:

```
removeEdge(g, (v,w)) :  
  Input graph g, edge (v,w)  
  
  if inLL(g.edges[v],w) then // (v,w) in graph  
    deleteLL(g.edges[v],w)  
    deleteLL(g.edges[w],v)  
    g.nE=g.nE-1  
  end if
```

Exercise #6:

Analyse storage cost and time complexity of adjacency list representation

Storage cost: $O(V+E)$

Cost of operations:

- initialisation: $O(V)$ (initialise V lists)
- insert edge: $O(1)$ (insert one vertex into list)
- delete edge: $O(E)$ (need to find vertex in list)

If vertex lists are sorted

- insert requires search of list $\Rightarrow O(E)$
- delete always requires a search, regardless of list order

Comparison of Graph Representations

	array of edges	adjacency matrix	adjacency list
space usage	E	V^2	$V+E$
initialise	1	V^2	V
insert edge	E	1	1
remove edge	E	1	E

Other operations:

	array of edges	adjacency matrix	adjacency list
disconnected(v)?	E	V	1
isPath(x,y)?	$E \cdot \log V$	V^2	$V+E$
copy graph	E	V^2	$V+E$
destroy graph	1	V	$V+E$

Graph Abstract Data Type

Graph ADT

Data:

- set of edges, set of vertices

Operations:

- building: create graph, add edge
- deleting: remove edge, drop whole graph
- scanning: check if graph contains a given edge

Things to note:

- set of vertices is fixed when graph initialised
- we treat vertices as **ints**, but could be arbitrary **Items**

Graph ADT (cont)

Graph ADT interface **graph.h**

```
// graph representation is hidden
typedef struct GraphRep *Graph;

// vertices denoted by integers 0..N-1
typedef int Vertex;

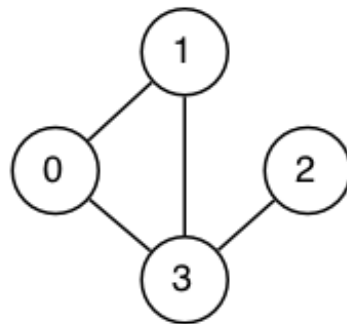
// edges are pairs of vertices (end-points)
typedef struct Edge { Vertex v; Vertex w; } Edge;

// operations on graphs
Graph newGraph(int V); // new graph with V vertices
void insertEdge(Graph, Edge);
void removeEdge(Graph, Edge);
bool adjacent(Graph, Vertex, Vertex); /* is there an edge
                                         between two vertices */
void freeGraph(Graph);
```


Graph ADT (Array of Edges)

Implementation of **GraphRep** (array-of-edges representation)

```
typedef struct GraphRep {
    Edge *edges; // array of edges
    int    nV;    // #vertices (numbered 0..nV-1)
    int    nE;    // #edges
    int    n;     // size of edge array
} GraphRep;
```



Undirected graph



Graph ADT (Array of Edges) (cont)

Implementation of graph initialisation (array-of-edges representation)

```
Graph newGraph(int V) {  
    assert(V >= 0);  
    Graph g = malloc(sizeof(GraphRep));    assert(g != NULL);  
  
    g->nV = V; g->nE = 0;  
    // allocate enough memory for edges  
    g->n = Enough;  
    g->edges = malloc(g->n*sizeof(Edge));    assert(g->edges != NULL);  
  
    return g;  
}
```

How much is enough? ... No more than $V(V-1)/2$... Much less in practice (sparse graph)

Graph ADT (Array of Edges) (cont)

Implementation of edge insertion/removal (array-of-edges representation)

```
// check if two edges are equal
bool eq(Edge e1, Edge e2) {
    return ( (e1.v == e2.v && e1.w == e2.w)
            || (e1.v == e2.w && e1.w == e2.v) );
}

void insertEdge(Graph g, Edge e) {
    // ensure that g exists and array of edges isn't full
    assert(g != NULL && g->nE < g->n);
    int i = 0;
    while (i < g->nE && !eq(e, g->edges[i]))
        i++;
    if (i == g->nE) // edge e not found
        g->edges[g->nE++] = e;
}

void removeEdge(Graph g, Edge e) {
    assert(g != NULL); // ensure that g exists
    int i = 0;
    while (i < g->nE && !eq(e, g->edges[i]))
        i++;
    if (i < g->nE) // edge e found
        g->edges[i] = g->edges[--g->nE];
}
```


Exercise #7: Checking Neighbours (i)

Assuming an array-of-edges representation ...

Implement a function to check whether two vertices are directly connected by an edge

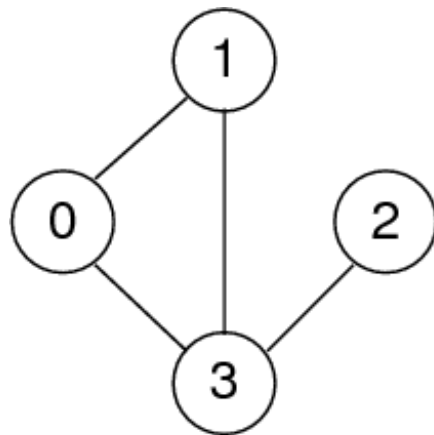
```
bool adjacent(Graph g, Vertex x, Vertex y) { ... }
```

```
bool adjacent(Graph g, Vertex x, Vertex y) {  
    assert(g != NULL);  
    Edge e;  
    e.v = x; e.w = y;  
    int i = 0;  
    while (i < g->nE) {  
        if (eq(e, g->edges[i])) // edge found  
            return true;  
        i++;  
    }  
    return false; // edge not found  
}
```

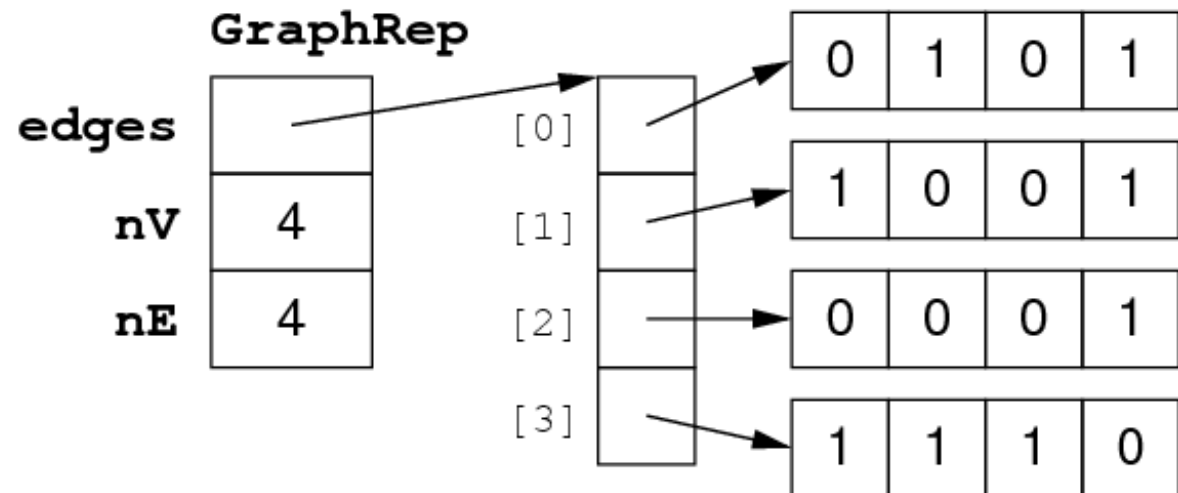
Graph ADT (Adjacency Matrix)

Implementation of **GraphRep** (adjacency-matrix representation)

```
typedef struct GraphRep {
    int    **edges; // adjacency matrix
    int     nV;     // #vertices
    int     nE;     // #edges
} GraphRep;
```



Undirected graph



Graph ADT (Adjacency Matrix) (cont)

Implementation of graph initialisation (adjacency-matrix representation)

```
Graph newGraph(int V) {
    assert(V >= 0);
    int i;

    Graph g = malloc(sizeof(GraphRep));          assert(g != NULL);
    g->nV = V;  g->nE = 0;

    // allocate memory for each row
    g->edges = malloc(V * sizeof(int *));        assert(g->edges != NULL);
    // allocate memory for each column and initialise with 0
    for (i = 0; i < V; i++) {
        g->edges[i] = calloc(V, sizeof(int));    assert(g->edges[i] != NULL);
    }
    return g;
}
```

standard library function **`calloc(size_t nelems, size_t nbytes)`**

- allocates a memory block of size **`nelems*nbytes`**
- and sets all bytes in that block to **zero**

Graph ADT (Adjacency Matrix) (cont)

Implementation of edge insertion/removal (adjacency-matrix representation)

```
// check if vertex is valid in a graph
bool validV(Graph g, Vertex v) {
    return (g != NULL && v >= 0 && v < g->nV);
}

void insertEdge(Graph g, Edge e) {
    assert(g != NULL && validV(g,e.v) && validV(g,e.w));

    if (!g->edges[e.v][e.w]) { // edge e not in graph
        g->edges[e.v][e.w] = 1;
        g->edges[e.w][e.v] = 1;
        g->nE++;
    }
}

void removeEdge(Graph g, Edge e) {
    assert(g != NULL && validV(g,e.v) && validV(g,e.w));

    if (g->edges[e.v][e.w]) { // edge e in graph
        g->edges[e.v][e.w] = 0;
        g->edges[e.w][e.v] = 0;
        g->nE--;
    }
}
```

Exercise #8: Checking Neighbours (ii)

Assuming an adjacency-matrix representation ...

Implement a function to check whether two vertices are directly connected by an edge

```
bool adjacent(Graph g, Vertex x, Vertex y) { ... }
```


```
bool adjacent(Graph g, Vertex x, Vertex y) {  
    assert(g != NULL && validV(g,x) && validV(g,y));  
  
    return (g->edges[x][y] != 0);  
}
```

Graph ADT (Adjacency List)

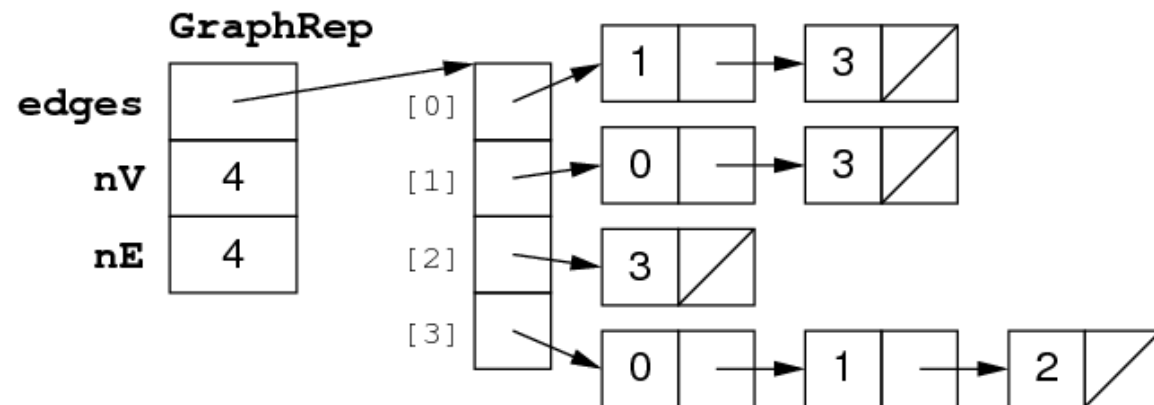
Implementation of **GraphRep** (adjacency-list representation)

```
typedef struct GraphRep {
    Node **edges;    // array of lists
    int    nV;       // #vertices
    int    nE;       // #edges
} GraphRep;

typedef struct Node {
    Vertex    v;
    struct Node *next;
} Node;
```



Undirected graph



Graph ADT (Adjacency List) (cont)

Implementation of graph initialisation (adjacency-list representation)

```
Graph newGraph(int V) {  
    assert(V >= 0);  
    int i;  
  
    Graph g = malloc(sizeof(GraphRep));          assert(g != NULL);  
    g->nV = V;  g->nE = 0;  
  
    // allocate memory for array of lists  
    g->edges = malloc(V * sizeof(Node *));  assert(g->edges != NULL);  
    for (i = 0; i < V; i++)  
        g->edges[i] = NULL;  
  
    return g;  
}
```

Graph ADT (Adjacency List) (cont)

Implementation of edge insertion/removal (adjacency-list representation)

```
void insertEdge(Graph g, Edge e) {
    assert(g != NULL && validV(g,e.v) && validV(g,e.w));

    if (!inLL(g->edges[e.v], e.w)) {        // edge e not in graph
        g->edges[e.v] = insertLL(g->edges[e.v], e.w);
        g->edges[e.w] = insertLL(g->edges[e.w], e.v);
        g->nE++;
    }
}

void removeEdge(Graph g, Edge e) {
    assert(g != NULL && validV(g,e.v) && validV(g,e.w));

    if (inLL(g->edges[e.v], e.w)) {        // edge e in graph
        g->edges[e.v] = deleteLL(g->edges[e.v], e.w);
        g->edges[e.w] = deleteLL(g->edges[e.w], e.v);
        g->nE--;
    }
}
```

inLL, **insertLL**, **deleteLL** are standard linked list operations (as discussed in week 3)

Exercise #9: Checking Neighbours (iii)

Assuming an adjacency list representation ...

Implement a function to check whether two vertices are directly connected by an edge

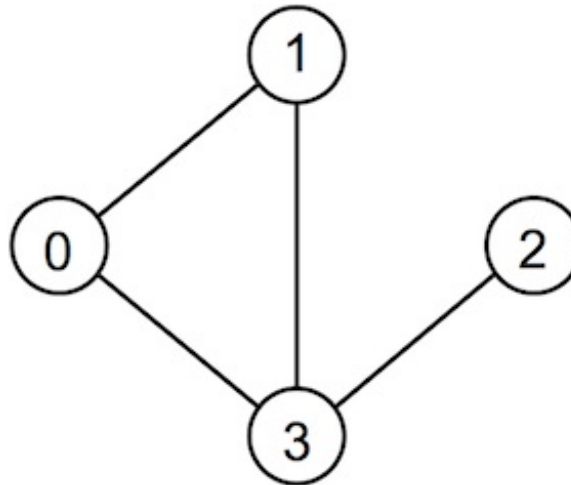
```
bool adjacent(Graph g, Vertex x, Vertex y) { ... }
```

```
bool adjacent(Graph g, Vertex x, Vertex y) {  
    assert(g != NULL && validV(g,x));  
  
    return inLL(g->edges[x], y);  
}
```

Exercise #10: Graph ADT Client

Write a program that uses the graph ADT to

- build the graph depicted below
- print all the nodes that are incident to vertex 1 in ascending order



```
#include <stdio.h>
#include "Graph.h"

#define NODES 4
#define NODE_OF_INTEREST 1

int main(void) {
    Graph g = newGraph(NODES);

    Edge e;
    e.v = 0; e.w = 1; insertEdge(g,e);
    e.v = 0; e.w = 3; insertEdge(g,e);
    e.v = 1; e.w = 3; insertEdge(g,e);
    e.v = 3; e.w = 2; insertEdge(g,e);

    int v;
    for (v = 0; v < NODES; v++) {
        if (adjacent(g, v, NODE_OF_INTEREST))
            printf("%d\n", v);
    }

    freeGraph(g);
    return 0;
}
```

Summary

- Graph terminology
 - vertices, edges, vertex degree, connected graph, tree
 - path, cycle, clique, spanning tree, spanning forest
- Graph representations
 - array of edges
 - adjacency matrix
 - adjacency lists
- Suggested reading:
 - Sedgewick, Ch.17.1-17.5