

Week 11

Things to Note ...

- Congratulations on finishing your second assignment!

In This Lecture ...

- Self-balancing trees ([S] Ch.13.1-13.4)

Coming Up ...

- Text processing algorithms

MyExperience Feedback

Assessment isn't a "one-way street" ...

- I get to assess you with my assignments and exams
- You get to assess me in UNSW's [MyExperience](#) Evaluation

Please fill it out ...

- give me some feedback on how you might like the course to run in the future
- even if that is "Exectaly the same. It was perfect this time."

Nerds You Should Know

Next in a series on famous computer scientists ...



She is considered to be the first computer programmer!

Nerds You Should Know (cont)

Ada Lovelace (1815-1852)



- Daughter of English poet Lord Byron
- Privately schooled in mathematics and science
- Met Charles Babbage as a teenager (1833)
- Published notes on ... (1843)
 - Babbage's **Analytical Engine**
 - algorithm for calculating "Bernoulli numbers" on AE
 - which thus became the first **computer program**
- Programming language **Ada** named after her (1980)

Nerds You Should Know (cont)

The first published computer algorithm ...

[illegible]

Tree Review

Binary search trees ...

- data structures designed for $O(\log n)$ search
- consist of nodes containing item (incl. key) and two links
- can be viewed as recursive data structure (subtrees)
- have overall ordering ($\text{data}(\text{Left}) < \text{root} < \text{data}(\text{Right})$)
- insert new nodes as leaves (or as root), delete from anywhere
- have structure determined by insertion order (worst: $O(n)$)
- operations: insert, delete, search, rotate, rebalance, ...

Randomised BST Insertion

Effects of order of insertion on BST shape:

- best case (for at-leaf insertion): keys inserted in pre-order (median key first, then median of lower half, median of upper half, etc.)
- worst case: keys inserted in ascending/descending order
- average case: keys inserted in **random** order $\Rightarrow O(\log_2 n)$

Tree ADT has no control over order that keys are supplied.

Can the algorithm itself introduce some **randomness**?

In the hope that this randomness helps to balance the tree ...

Sidetrack: Random Numbers

How can a computer pick a number at random?

- it cannot

Software can only produce **pseudo random numbers**.

- a pseudo random number is one that is predictable
 - (although it may appear unpredictable)

⇒ Implementation may deviate from expected theoretical behaviour

Sidetrack: Random Numbers (cont)

The most widely-used technique is called the **Linear Congruential Generator (LCG)**

- it uses a **recurrence** relation:
 - $X_{n+1} = (a \cdot X_n + c) \bmod m$, where:
 - m is the "modulus"
 - a , $0 < a < m$ is the "multiplier"
 - c , $0 \leq c \leq m$ is the "increment"
 - X_0 is the "**seed**"
 - if $c=0$ it is called a **multiplicative congruential generator**

LCG is not good for applications that need extremely high-quality random numbers

- the period length is too short (length of the sequence at which point it repeats itself)
- a short period means the numbers are correlated

Sidetrack: Random Numbers (cont)

Trivial example:

- for simplicity assume $c=0$
- so the formula is $X_{n+1} = a \cdot X_n \bmod m$
- try $a=11=X_0$, $m=31$, which generates the sequence:

```
11, 28, 29, 9, 6, 4, 13, 19, 23, 5, 24, 16, 21, 14, 30, 20, 3, 2, 22, 25,
27, 18, 12, 8, 26, 7, 15, 10, 17, 1, 11, 28, 29, 9, 6, 4, 13, 19, 23, 5,
24, 16, 21, 14, 30, 20, 3, 2, 22, 25, 27, 18, 12, 8, 26, 7, 15, 10, 17, 1,
11, 28, 29, 9, 6, 4, 13, 19, 23, 5, 24, 16, 21, 14, 30, 20, 3, 2, 22, 25, 27,
18, 12, 8, 26, 7, 15, 10, 17, 1, 11, 28, 29, 9, 6, 4, 13, 19, 23, 5, 24, 16,
21, 14, 30, 20, 3, 2, 22, 25, 27, 18, 12, 8, 26, 7, 15, 10, 17, 1, 11, 28,
29, 9, 6, 4, 13, 19, 23, 5, 24, 16, 21, 14, 30, 20, 3, 2, 22, 25, 27, 18,
12, 8, 26, 7, 15, 10, 17, 1, ...
```

- all the integers from 1 to 30 are here

Sidetrack: Random Numbers (cont)

Another trivial example:

- again let $c=0$
- try $a=12=X_0$ and $m=30$
 - that is, $X_{n+1} = 12 \cdot X_n \bmod 30$
 - which generates the sequence:

12, 24, 18, 6, 12, 24, 18, 6, 12, 24, 18, 6, 12, 24, 18, 6, 12, 24, 18, 6,
12, 24, 18, 6, 12, 24, 18, 6, 12, 24, 18, 6, ...

- notice the period length ... clearly a terrible sequence

Sidetrack: Random Numbers (cont)

It is a complex task to pick good numbers. A bit of history:

Lewis, Goodman and Miller (1969) suggested

- $X_{n+1} = 7^5 \cdot X_n \bmod (2^{31}-1)$
- note:
 - 7^5 is 16807
 - $2^{31}-1$ is 2147483647
 - $X_0 = 0$ is not a good seed value

Most compilers use LCG-based algorithms that are slightly more involved; see www.mscs.dal.ca/~selinger/random/ for details (including a short C program that produces the exact same pseudo-random numbers as **gcc** for any given seed value)

Sidetrack: Random Numbers (cont)

- Two functions are required:

```
srandom(int seed) // sets its argument as the seed
```

```
random() // uses a LCG technique to generate random  
         // numbers in the range 0 .. RAND_MAX
```

where the constant **RAND_MAX** is defined in **stdlib.h**
(depends on the computer: on the CSE network, RAND_MAX = 2147483647)

- The period length of this random number generator is very large
approximately $16 \cdot ((2^{31}) - 1)$

Sidetrack: Random Numbers (cont)

To convert the return value of `random()` to a number between 0 .. RANGE

- compute the remainder after division by `RANGE+1`

Using the remainder to compute a random number is not the best way:

- can generate a 'better' random number by using a more complex division
- but good enough for most purposes

Some applications require more sophisticated, *cryptographically secure* pseudo random numbers

Exercise #1: Random Numbers

Write a program to simulate 10,000 rounds of Two-up.

- Assume a \$10 bet at each round
- Compute the overall outcome and average per round

```
#include <stdlib.h>
#include <stdio.h>

#define RUNS 10000
#define BET 10

int main(void) {
    int coin1, coin2, n, sum = 0;
    for (n = 0; n < RUNS; n++) {
        do {
            coin1 = random() % 2;
            coin2 = random() % 2;
        } while (coin1 != coin2);
        if (coin1==1 && coin2==1)
            sum += BET;
        else
            sum -= BET;
    }
    printf("Final result: %d\n", sum);
    printf("Average outcome: %f\n", (float) sum / RUNS);
    return 0;
}
```


Sidetrack: Random Numbers (cont)

Seeding

There is one significant problem:

- every time you run a program with the same seed, you get exactly the same sequence of 'random' numbers (why?)

To vary the output, can give the random seeder a starting point that varies with time

- an example of such a starting point is the current time, `time(NULL)`
(NB: this is different from the UNIX command `time`, used to measure program running time)

```
#include <time.h>
time(NULL) // returns the time as the number of seconds
           // since the Epoch, 1970-01-01 00:00:00 +0000

// time(NULL) on October 10th, 2017, 12:59pm was 1507600763
// time(NULL) about a minute later was 1507600825
```

Randomised BST Insertion

Approach: normally do leaf insert, randomly do root insert.

```
insertRandom(tree,item)
  Input  tree, item
  Output tree with item randomly inserted

  if tree is empty then
    return new node containing item
  end if
  // p/q chance of doing root insert
  if random() mod q < p then
    return insertAtRoot(tree,item)
  else
    return insertAtLeaf(tree,item)
  end if
```

E.g. 30% chance \Rightarrow choose $p=3, q=10$

Randomised BST Insertion (cont)

Cost analysis:

- similar to cost for inserting keys in random order: $O(\log_2 n)$
- does not rely on keys being supplied in random order

Approach can also be applied to deletion:

- standard method promotes inorder successor to root
- for the randomised method ...
 - promote inorder successor from right subtree, OR
 - promote inorder predecessor from left subtree

Splay Trees

Splay Trees

A kind of "self-balancing" tree ...

Splay tree insertion modifies insertion-at-root method:

- by considering **p**arent-**c**hild-**g**ranchild (three level analysis)
- by performing double-rotations based on p-c-g orientation

The idea: appropriate double-rotations improve tree balance.


Splay tree implementations also do **rotation-in-search**:

- can provide similar effect to periodic rebalance
- improves balance, but makes search more expensive

Splay Trees (cont)


Cases for splay tree double-rotations:

- case 1: grandchild is left-child of left-child
- case 2: grandchild is right-child of left-child
- case 3: grandchild is left-child of right-child
- case 4: grandchild is right-child of right-child




Splay Trees (cont)

Example: double-rotation case for left-child of left-child:



Splay Trees (cont)

Example: double-rotation case for right-child of left-child:



Splay Trees (cont)

Algorithm for splay tree insertion:

```


insertSplay(tree,item):
  Input  tree, item
  Output tree with item splay-inserted

  if tree is empty then return new node containing item
  else if item=data(tree) then return tree
  else if item<data(tree) then
    if left(tree) is empty then
      left(tree)=new node containing item
    else if item<data(left(tree)) then
      // Case 1: left-child of left-child
      left(left(tree))=insertSplay(left(left(tree)),item)
      left(tree)=rotateRight(left(tree))
    else // Case 2: right-child of left-child
      right(left(tree))=insertSplay(right(left(tree)),item)
      left(tree)=rotateLeft(left(tree))
    end if
    return rotateRight(tree)
  else if item>data(tree) then
    if right(tree) is empty then
      right(tree)=new node containing item
    else if item<data(right(tree)) then
      // Case 3: left-child of right-child
      left(right(tree))=insertSplay(left(right(tree)),item)
      right(tree)=rotateRight(right(tree))
    else // Case 4: right-child of right-child
      right(right(tree))=insertSplay(right(right(tree)),item)
      right(tree)=rotateLeft(right(tree))
    end if
    return rotateLeft(tree)
  end if


```

Exercise #2: Splay Trees

Insert **18** into this splay tree:



Correction:



Splay Trees (cont)

Searching in splay trees:


```
searchSplay(tree,item):
  Input   tree, item
  Output address of item if found in tree
           NULL otherwise

  if tree=NULL then
    return NULL
  else
    tree=splay(tree,item)
    if data(tree)=item then
      return tree
    else
      return NULL
    end if
  end if
```


where **splay()** is similar to **insertSplay()**,
except that it doesn't add a node ... simply moves **item** to root if found, or
nearest node if not found

Exercise #3: Splay Trees

If we search for **22** in the splay tree



... how does this affect the tree?



Splay Trees (cont)

Analysis of splay tree performance:

- assume that we "splay" for both insert and search
- consider: m insert+search operations, n nodes
- total number of comparisons: average $O((n+m) \cdot \log_2(n+m))$

Gives good overall (amortized) cost.

- insert cost not significantly different to insert-at-root
- search cost increases, but ...
 - improves balance on each search
 - moves frequently accessed nodes closer to root

But ... still has worst-case search cost $O(n)$

Real Balanced Trees

Better Balanced Binary Search Trees

So far, we have seen ...

- randomised trees ... make poor performance unlikely
- occasional rebalance ... fix balance periodically
- splay trees ... reasonable amortized performance
- but both types still have $O(n)$ worst case

Ideally, we want both average/worst case to be $O(\log n)$

- AVL trees ... fix imbalances as soon as they occur
- 2-3-4 trees ... use varying-sized nodes to assist balance
- red-black trees ... isomorphic to 2-3-4, but binary nodes

AVL Trees

AVL Trees

Invented by Georgy Adelson-Velsky and Evgenii Landis

Approach:

- insertion (at leaves) may cause imbalance
- repair balance as soon as we notice imbalance
- repairs done locally, not by overall tree restructure

A tree is unbalanced when: $\text{abs}(\text{height}(\text{left}) - \text{height}(\text{right})) > 1$

This can be repaired by a single rotation:

- if left subtree too deep, rotate right
- if right subtree too deep, rotate left

Problem: determining height/depth of subtrees may be expensive.

AVL Trees (cont)

Implementation of AVL insertion

```


insertAVL(tree,item):
    Input  tree, item
    Output tree with item AVL-inserted


    if tree is empty then
        return new node containing item
    else if item=data(tree) then
        return tree
    else
        if item<data(tree) then
            left(tree)=insertAVL(left(tree),item)
        else if item>data(tree) then
            right(tree)=insertAVL(right(tree),item)
        end if
        if height(left(tree))-height(right(tree)) > 1 then
            tree=rotateRight(tree)
        else if height(right(tree))-height(left(tree)) > 1 then
            tree=rotateLeft(tree)
        end if
        return tree
    end if

```

Exercise #4: AVL Trees

Insert **27** into the AVL tree






AVL Trees (cont)

Analysis of AVL trees:

- trees are **height**-balanced; subtree depths differ by ± 1
- average/worst-case search performance of $O(\log n)$
- *require* extra data to be stored in each node (efficiency)
- may not be **weight**-balanced; subtree sizes may differ




2-3-4 Trees

2-3-4 Trees


2-3-4 trees have three kinds of nodes

- 2-nodes, with two children (same as normal BSTs)
- 3-nodes, two values and three children
- 4-nodes, three values and four children



2-3-4 Trees (cont)

2-3-4 trees are ordered similarly to BSTs



In a **balanced 2-3-4 tree**:

- all leaves are at same distance from the root

2-3-4 trees grow "upwards" from the leaves.

2-3-4 Trees (cont)

Possible 2-3-4 tree data structure:

```
typedef struct node {  
    int          order;          // 2, 3 or 4  
    int          data[3];        // items in node  
    struct node *child[4];       // links to subtrees  
} node;
```

2-3-4 Trees (cont)

Searching in 2-3-4 trees:

```
Search(tree,item):
    Input   tree, item
    Output address of item if found in 2-3-4 tree
            NULL otherwise

    if tree is empty then
        return NULL
    else
        i=0
        while i<tree.order-1 ^ item>tree.data[i] do
            i=i+1    // find relevant slot in data[]
        end while
        if item=tree.data[i] then    // item found
            return address of tree.data[i]
        else    // keep looking in relevant subtree
            return Search(tree.child[i],item)
        end if
    end if
```

2-3-4 Trees (cont)

2-3-4 tree searching cost analysis:

- as for other trees, worst case determined by height h
- 2-3-4 trees are always balanced \Rightarrow height is $O(\log n)$
- worst case for height: all nodes are 2-nodes
same case as for balanced BSTs, i.e. $h \approx \log_2 n$
- best case for height: all nodes are 4-nodes
balanced tree with branching factor 4, i.e. $h \approx \log_4 n$



Insertion into 2-3-4 Trees

Insertion algorithm:


- find leaf node where Item belongs (via search)
- if not full (i.e. order < 4)
 - insert Item in this node, order++
- if node is full (i.e. contains 3 items)
 - split into two 2-nodes as leaves
 - promote middle element to parent
 - insert item into appropriate leaf 2-node
 - if parent is a 4-node
 - continue split/promote upwards
 - if promote to root, and root is a 4-node
 - split root node and add new root

Insertion into 2-3-4 Trees (cont)

Insertion into a 2-node or 3-node:




Insertion into a 4-node (requires a split):




Insertion into 2-3-4 Trees (cont)

Splitting the root:




Insertion into 2-3-4 Trees (cont)


Building a 2-3-4 tree ... 7 insertions:



Exercise #5: Insertion into 2-3-4 Tree

Show what happens when D, S, F, U are inserted into this tree:





Insertion into 2-3-4 Trees (cont)

Insertion algorithm:

```

insert(tree,item):
    Input  2-3-4 tree, item
    Output tree with item inserted

    if tree is empty then
        return new node containing item
    end if
    node=Search(tree,item)
    parent=parent of node
    if node.order<4 then
        insert item into node
        increment node.order
    else
        promote = node.data[1]           // middle value
        nodeL   = new node containing data[0]
        nodeR   = new node containing data[2]
        if item<node.data[1] then
            insert(nodeL,item)
        else
            insert(nodeR,item)
        end if
        insert(parent,promote)
        while parent.order=4 do
            continue promote/split upwards
        end while
        if parent is root ^ parent.order=4 then
            split root, making new root
        end if
    end if
end if

```

Insertion into 2-3-4 Trees (cont)

Variations on 2-3-4 trees ...

Variation #1: why stop at 4? why not 2-3-4-5 trees? or M -way trees?

- allow nodes to hold up to $M-1$ items, and at least $M/2$
- if each node is a disk-page, then we have a **B-tree** (databases)
- for B-trees, depending on **Item** size, $M > 100/200/400$

Variation #2: don't have "variable-sized" nodes

- use standard BST nodes, augmented with one extra piece of data
- implement similar strategy as 2-3-4 trees → red-black trees.

Red-Black Trees

Red-Black Trees

Red-black trees are a representation of 2-3-4 trees using BST nodes.

- each node needs one extra value to encode link type
- but we no longer have to deal with different kinds of nodes

Link types:

- red links ... combine nodes to represent 3- and 4-nodes
- black links ... analogous to "ordinary" BST links (child links)

Advantages:

- standard BST search procedure works unmodified
- get benefits of 2-3-4 tree self-balancing (although deeper)

Red-Black Trees

Definition of a red-black tree

- a BST in which each node is marked red or black
- no two red nodes appear consecutively on any path
- a red node corresponds to a 2-3-4 sibling of its parent
- a black node corresponds to a 2-3-4 child of its parent

Balanced red-black tree


- all paths from root to leaf have same number of black nodes

Insertion algorithm: avoids worst case $O(n)$ behaviour

Search algorithm: standard BST search

Red-Black Trees (cont)

Representing 4-nodes in red-black trees:




Some texts colour the links rather than the nodes.


Red-Black Trees (cont)

Representing 3-nodes in red-black trees (two possibilities):




2-3-4 nodes



red-black nodes (i)




red-black nodes (ii)



Red-Black Trees (cont)

Equivalent trees (one 2-3-4, one red-black):



Red-Black Trees (cont)

Red-black tree implementation:

```
typedef enum {RED, BLACK} Colour;
typedef struct node *RBTree;
typedef struct node {
    int      data;          // actual data
    Colour colour;          // relationship to parent
    RBTree  left;           // left subtree
    RBTree  right;          // right subtree
} node;

#define colour(tree) ((tree)->colour)
#define isRed(tree)  ((tree) != NULL && (tree)->colour == RED)
```

RED = node is part of the same 2-3-4 node as its parent (sibling)

BLACK = node is a child of the 2-3-4 node containing the parent

Red-Black Trees (cont)


New nodes are always **red**:

```
RBTree newNode(Item it) {  
    RBTree new = malloc(sizeof(Node));  
    assert(new != NULL);  
    data(new) = it;  
    colour(new) = RED;  
    left(new) = right(new) = NULL;  
    return new;  
}
```

Red-Black Trees (cont)

Node.red allows us to distinguish links

- black = parent node is a "real" parent
- red = parent node is a 2-3-4 neighbour



Red-Black Trees (cont)

Search method is standard BST search:

```
SearchRedBlack(tree,item):  
    Input   tree, item  
    Output true if item found in red-black tree  
             false otherwise  
  
    if tree is empty then  
        return false  
    else if item < data(tree) then  
        return SearchRedBlack(left(tree),item)  
    else if item > data(tree) then  
        return SearchRedBlack(right(tree),item)  
    else                                     // found  
        return true  
    end if
```

Red-Black Tree Insertion

Insertion is more complex than for standard BSTs

- need to recall direction of last branch (L or R)
- need to recall whether parent link is red or black
- splitting/promoting implemented by rotateLeft/rotateRight
- several cases to consider depending on colour/direction combinations

Red-Black Tree Insertion (cont)

High-level description of insertion algorithm:

```
insertRB(tree,item,inRight):
```

Input tree, item, inRight indicating direction of last branch

Output tree with it inserted

if tree is empty **then**

return newNode(item)

end if

if left(tree) and right(tree) both are **RED** **then**

split 4-node

end if

recursive insert cases (cf. regular BST)

re-arrange links/colours after insert

return modified tree

```
insertRedBlack(tree,item):
```

Input red-black tree, item

Output tree with item inserted


tree=insertRB(tree,item,false)

colour(tree)=**BLACK**

return tree

Red-Black Tree Insertion (cont)

Splitting a 4-node, in a red-black tree:



Algorithm:

```

if isRed(left(currentTree)) ^ isRed(right(currentTree)) then
    colour(currentTree)=RED
    colour(left(currentTree))=BLACK
    colour(right(currentTree))=BLACK
end if
    
```

Red-Black Tree Insertion (cont)

Simple recursive insert (a la BST):




Algorithm:

```
if item < data(tree) then
    left(tree) = insertRB(left(tree), item, false)
    re-arrange links/colours after insert
else // item larger than data in root
    right(tree) = insertRB(right(tree), item, true)
    re-arrange links/colours after insert
end if
```

Not affected by colour of **tree** node.

Red-Black Tree Insertion (cont)

Re-arrange after insert (1): two successive red links = newly-created 4-node




Algorithm:

```
if isRed(left(currentTree)) ^ isRed(left(left(currentTree))) then
  currentTree=rotateRight(currentTree)
  colour(currentTree)=BLACK
  colour(right(currentTree))=RED
end if
```

Red-Black Tree Insertion (cont)

Re-arrange after insert (2): "normalise" direction of successive red links




Algorithm:

```
if inRight  $\wedge$  isRed(currentTree)  $\wedge$  isRed(left(currentTree)) then
    currentTree=rotateRight(currentTree)
end if
```

Red-Black Tree Insertion (cont)

Example of insertion, starting from empty tree:

22, 12, 8, 15, 11, 19, 43, 44, 45, 42, 41, 40, 39



Red-black Tree Performance

Cost analysis for red-black trees:

- tree is well-balanced; worst case search is $O(\log_2 n)$
- insertion affects nodes down one path; max #rotations is $2 \cdot h$
(where h is the height of the tree)

Only disadvantage is complexity of insertion/deletion code.

Note: red-black trees were popularised by Sedgewick.

Summary

- Random numbers
- Self-adjusting trees
 - Splay trees
 - AVL trees
 - 2-3-4 trees
 - Red-black trees
- Suggested reading:
 - Sedgewick, Ch.13.1-13.4